

# *Efficient Visual Understanding and Interaction with VLMs*

Wentong LI (李文通)

Homepage: <https://cslwt.github.io/>

Associate Professor@NUAA; PhD@ZJU

2025/08/09

# Content

1. Fine-grained Object & Region Understanding
  - Image/Video
2. Efficient VLMs via Visual Token Compression
  - Model-driven/Data-driven
3. Streaming Understanding & Interaction for AI Assistant
  - Training/Training-free

# **Content**

- 1. Fine-grained Object & Region Understanding**
  - Image/Video**
- 2. Efficient VLMs via Visual Token Compression**
  - Model-driven/Data-driven**

# Fine-grained Object/Region Understanding

## Osprey



SAM "Segment Everything" Predictions

**No semantic information**

- Integrate images, target regions (masks), and textural data;
- Enable fine-grained semantic description of arbitrary regions or objects within images;
- Strong robustness and generalization.

Object Category: person

Part Taxonomy: body

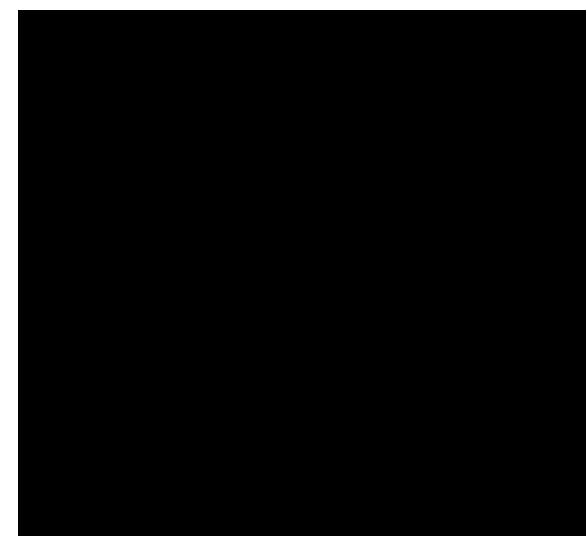
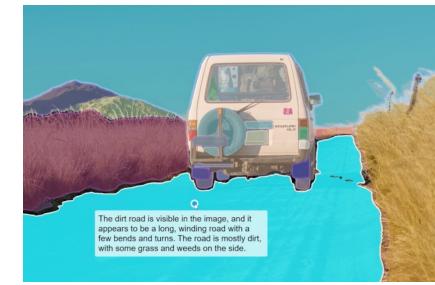
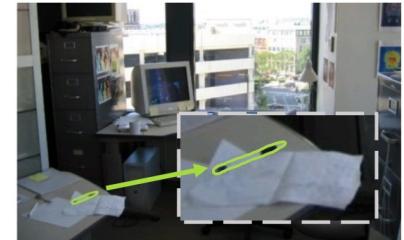
Attribute: color, position ...

Caption: region short / detailed description

Fine-grained Region/Pixel Understanding

**Rich semantic information containing different granularities**

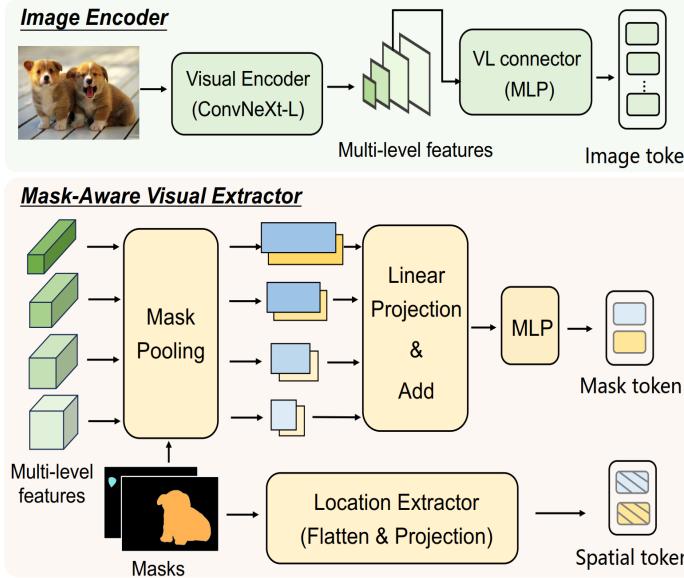
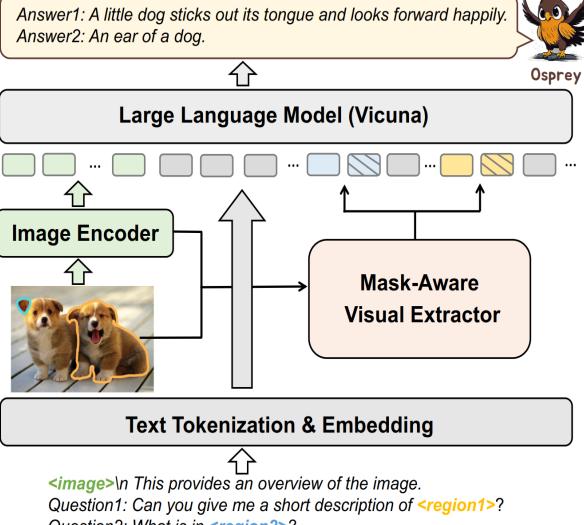
General scene



Out-of-domain Scene

# Fine-grained Object/Region Understanding

## Osprey



Object Categories	
①	dog
②	dog collar
③	tag
④	chair
⑤	cushion
⑥	dinner table
⑦	cookie
⑧	cake
⑨	box

- Part Attributes**
- ⑪ Q: What is the category?  
A: Side of a box.  
Q: What is the material?  
A: Paper.
  - ⑫ Q: Is the inner side of the box plain in <region12> ?  
A: Yes.  
Q: What is the color?  
A: White.



## Detailed Descriptions

- ① The chair behind the dog has a brown color and appears to be made of wood. It has a simple design that complements the cushions' blue color on which the dog is resting.
- ② In front of the dog and on the dining table, there are several decorative dog treats and cakes. One of them looks like a big white bone, while others have dots on them. There is also a dotted item in front of the treats on the left side.

## Robustness and Flexibility

- ③ Q: Cat is the category of <region1>, right?  
A: No, there is no cat in the region.
- ④ Q: Is the category of <region2> dog?  
A: No, it isn't.
- ⑤ Q: Where is the wood part of the chair located in <region4>?  
Answer the question using a single word or phrase.  
A: Behind the dog.

## Conversations and Complex Reasoning

- Question: Can you describe the pillow that the dog is resting on in <region5>?  
Answer: The pillow in <region5> is a blue cushion or pillow that is part of the chair. It provides a comfortable resting place for the dog. The cushion is visible in the bottom part of the image and supports the dog's brown paws as it rests on the chair.
- Question: What can you tell me about the items on the cake in <region7>, <region8> and <region9>?  
Answer: In all regions, there are treats or cakes displayed. In <region9>, there is a plated item with dots or a pattern on it, placed above a big white bone. It seems to be a celebratory arrangement of doggy treats and cakes. Meanwhile, in <region7> and <region8>, there are two small dog treats or bones sitting on a plate. These treats may be specially made for dogs with some writing or decoration on them.
- Question: Can you provide any more details about the dog and its surroundings based on the descriptions?  
Answer: The dog in the image is black and brown and is sitting on a blue chair cushion, as seen in <region5>. The dog is surrounded by celebratory doggy treats and cakes, as mentioned in <region7> and <region8>. The presence of the treats and the dog's position on the chair suggests a special occasion or celebration. The overall theme of the image seems to be centered around the dog and its enjoyment of the treats and cakes.

- Support high-resolution image
  - ConvNeXt (512x521@training, 800x800@inference)
- Pixel-level region feature extraction
  - Mask-Aware visual extractor (multi-level)

- 720K region-text pairs.
- Six types of object region-text data.

Open-source: <https://huggingface.co/sunshine-lwt>



# Fine-grained Object/Region Understanding

## Online Demo

Segmentation



point-prompt box-prompt segment everything

Input image



Description(Osprey)

red light on top of police car.

Generate segmentation and description

Clear Image

Segmentation

Description(Osprey)



Video Visualization

# Fine-grained Object/Region Understanding

## Spatial RGPT

Accurate Spatial Understanding & Reasoning on Objects in Image.

- Relative Relation
- Metric Measurement
- Spatial Reasoning



How wide is Region [1]?

The width of Region [1] is 7.73 feet.

Is Region [1] behind Region [4]?

No, it is in front of Region [4].

Between Region [0] and Region [2], which one has more height?

Region [2] is taller.

### A Relative Relation

⑧ Can you confirm if ④ has a greater width compared to ⑥?

⑨ In fact, ④ might be narrower than ⑥.

⑧ Which of these four, ③, ④, ⑤, ⑥ is the tallest?

⑨ ⑥ is the tallest.

⑧ From the viewer's perspective, is the position of ⑥ closer than that of ③?

⑨ No, ③ is closer.

### B Metric Measurement

⑧ If you are at ③, where will you find ⑦?

⑨ ⑦ is at 1 o'clock from ③.

⑧ What is the height of ④?

⑨ ④ is 1.38 meters tall.

⑧ How far apart are ② and ③ horizontally?

⑨ They are 15.33 meters apart horizontally.

### C Spatial Reasoning

⑧ Can you try to estimate how many floors are in ①? Think step by step.

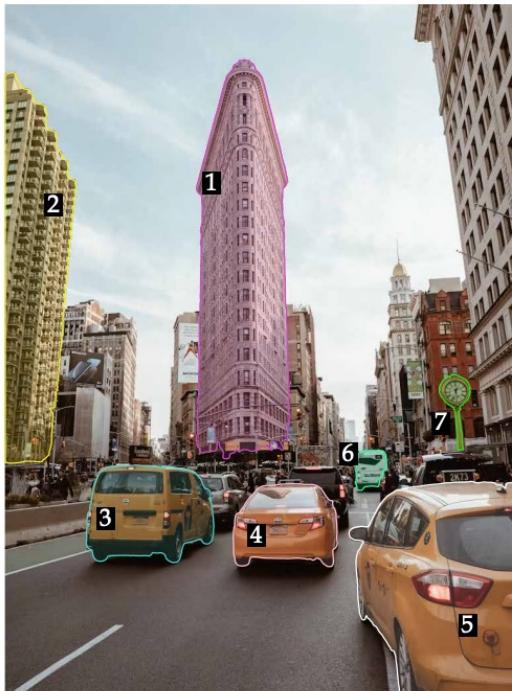
⑨ The height of ① is 204.54 feet. Assuming each floor is about 10 feet high, the total number of floors would be 20.454. Since you can't have a fraction of a floor, the total number of floors would be approximately 20.

⑧ I'm planning to build a bridge from ② to ①, is it possible?

⑨ The distance between ② and ① is 14.82 meters, so it is possible to build a bridge between them.

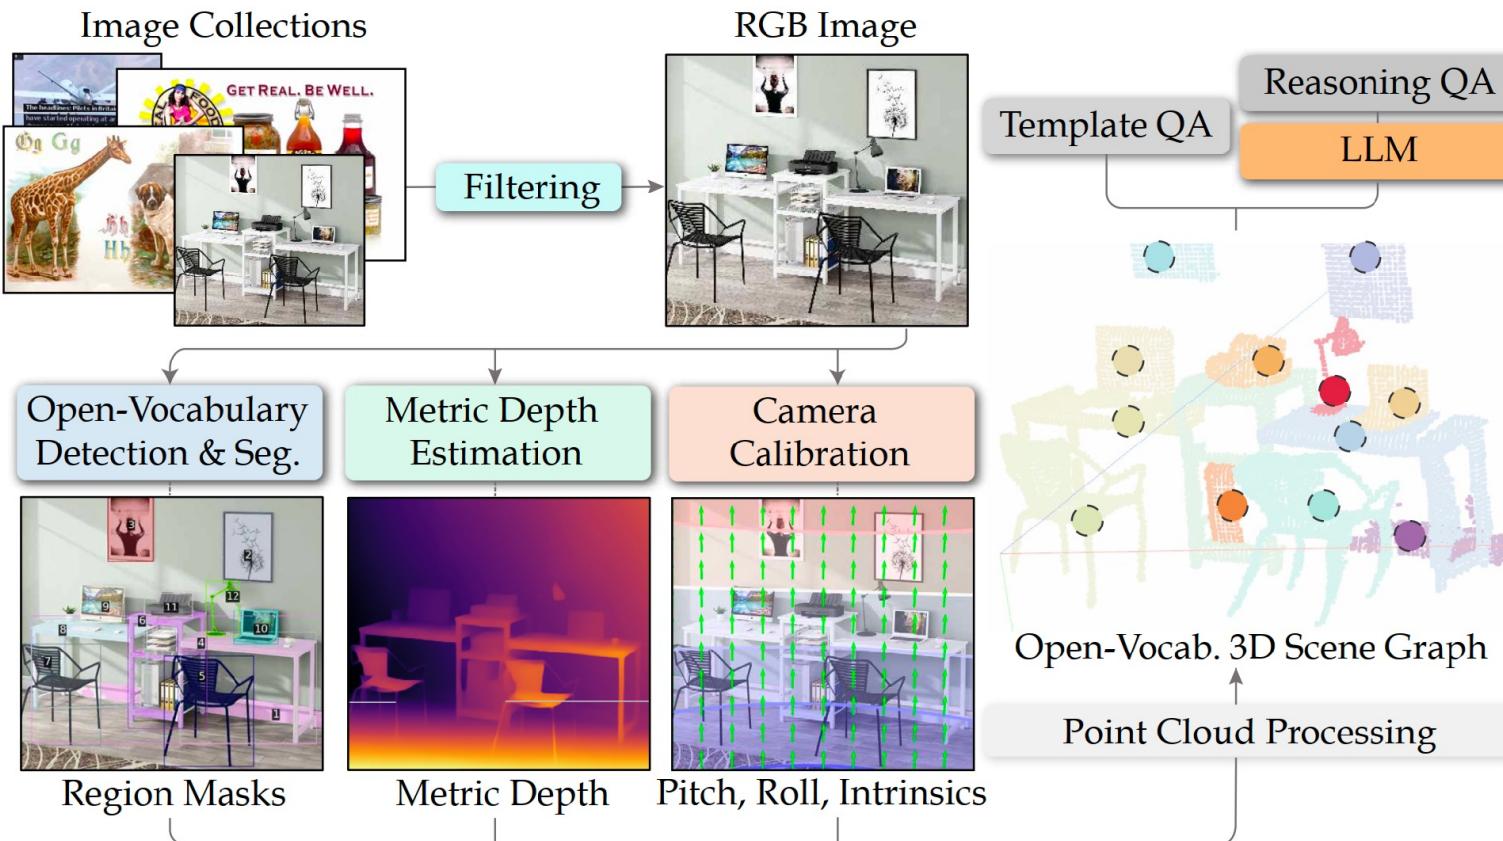
⑧ If you are riding a motorcycle with 36 inches wide, do you think you can pass through the area between ③ and ④?

⑨ The distance between ③ and ④ is 38.95 inches, so yes, you can pass through the area between ③ and ④ since the motorcycle is narrower than the distance between them.



# Fine-grained Object/Region Understanding

## 3D scene graph construction from Single 2D image



- Open-Voc. Detection & Segmentation:  
Tagging model, Grounding DINO, SAM-HQ
- Metric Depth Estimation:  
Metric3Dv2
- Camera Calibration:  
WildCamera: camera intrinsic  
PerspectiveFields: camera extrinsics
- 3D Scene Graph Construction

# Fine-grained Object/Region Understanding

## Framework

*They are around 45 centimeters apart.*

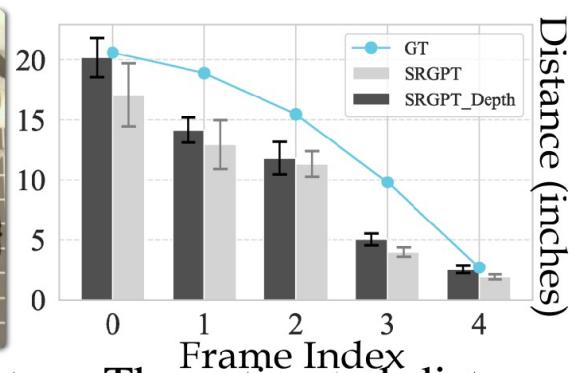
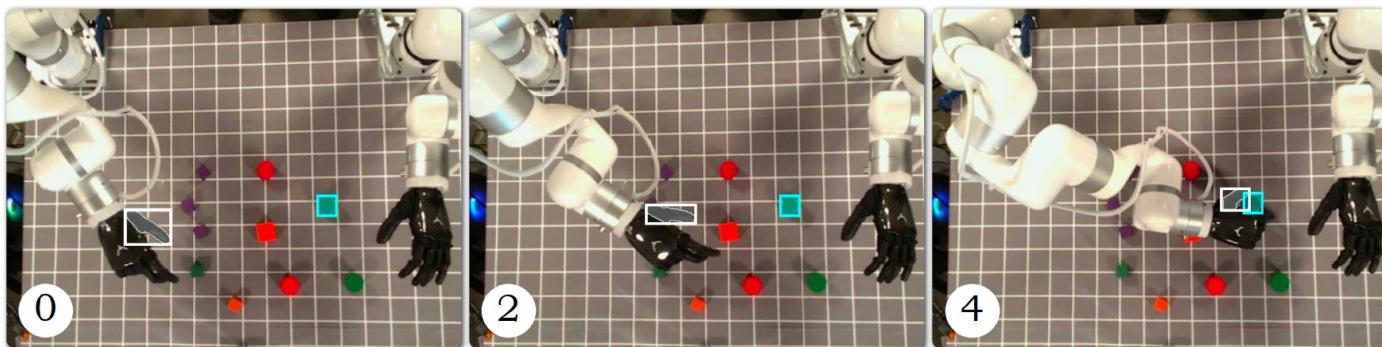
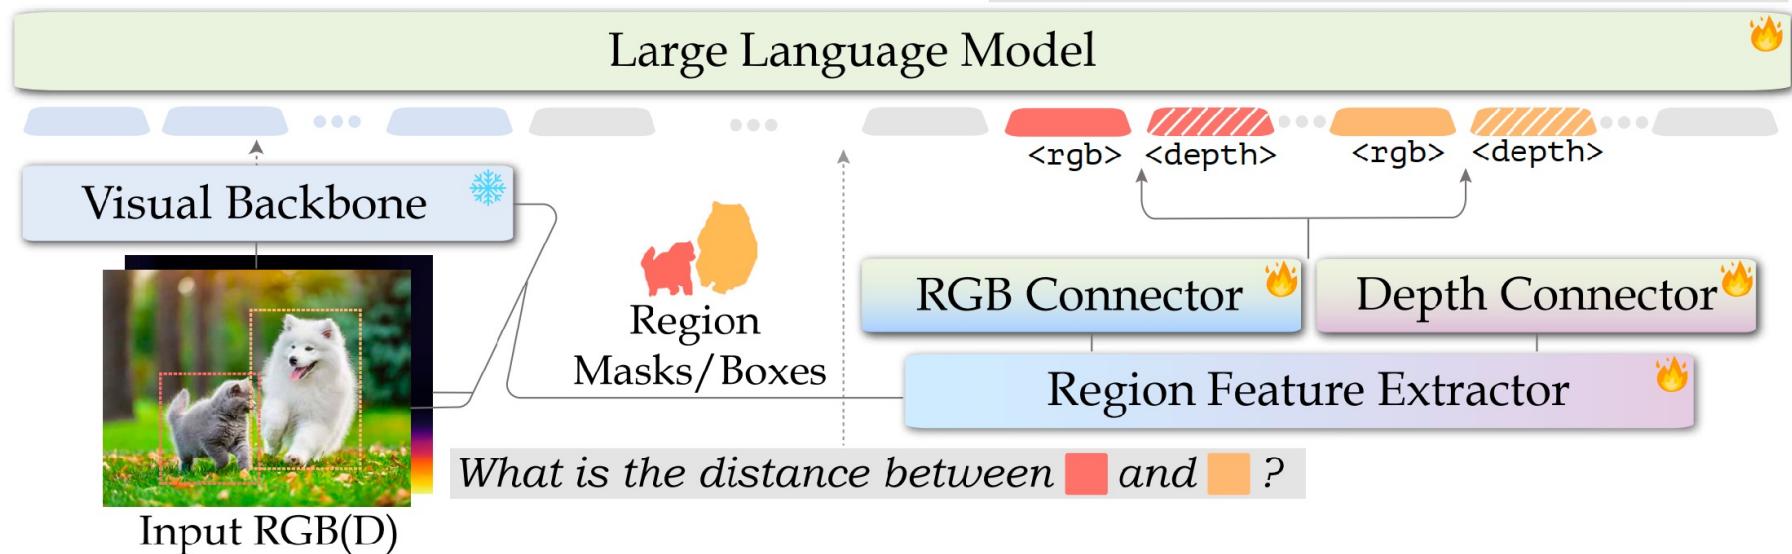
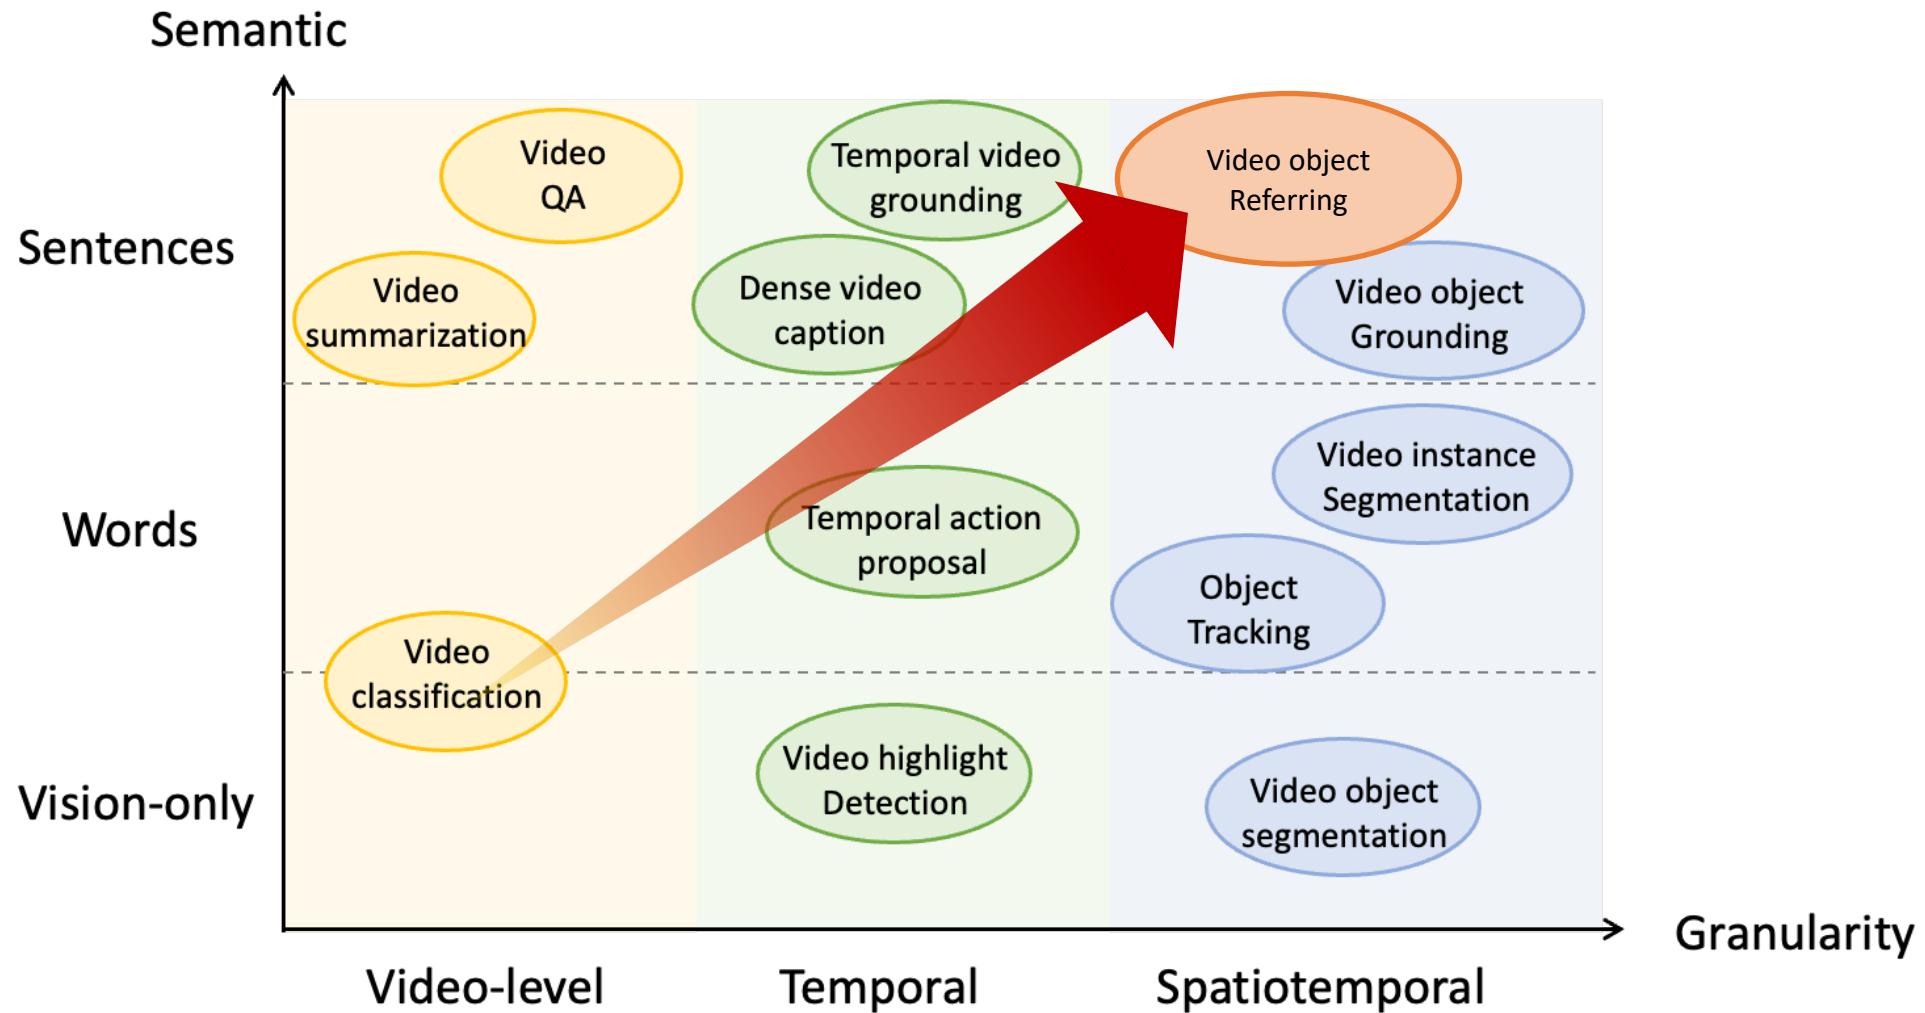


Figure 6: SpatialRGPT functions as a region-aware reward annotator. The estimated distance decreased monotonically as the fingertip moves towards the target.

# Fine-grained Object/Region Understanding



# Fine-grained Object/Region Understanding

## Video Object Referring



A man with a cocked hat and green robes, riding a horse, slowly riding from the left to the right.

## Future Reasoning



Q: What will <object1> probably do next?

A: <object1> will probably have to shoot or pass the ball to a teammate.

## Video Objects Relationship



The knife <object1> moves the spring onions from the chopping board <object2> to the pan.

## Input image



## Video Object Retrieval



The man was Trump, who stood in the crowd waving and waving his fist to the left and right.

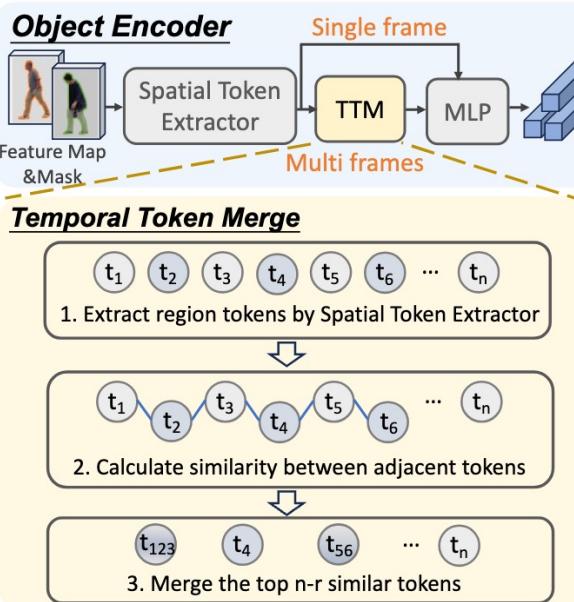
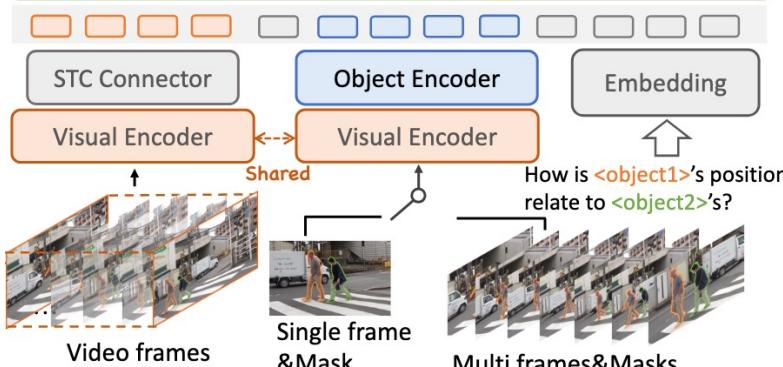
# Fine-grained Object/Region Understanding

## VideoRefer Suite

object1 and object2 are walking in the same direction, with object2 consistently trailing behind object1 throughout the entire sequence.

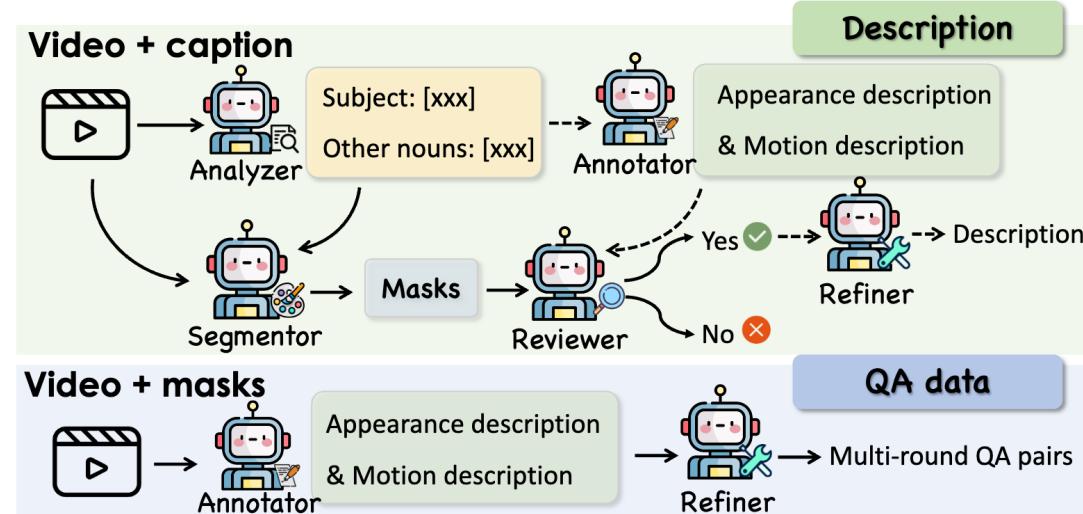


### Large Language Model

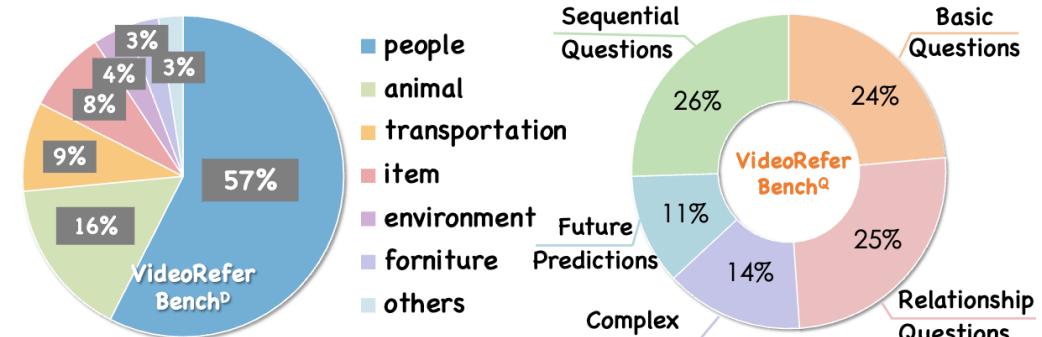


### VideoRefer Model

- Spatiotemporal Region-level understanding Architecture;
- Constructing Large-scale Video Region Dataset;
- Evaluation Benchmarks for Video-based Object Understanding.



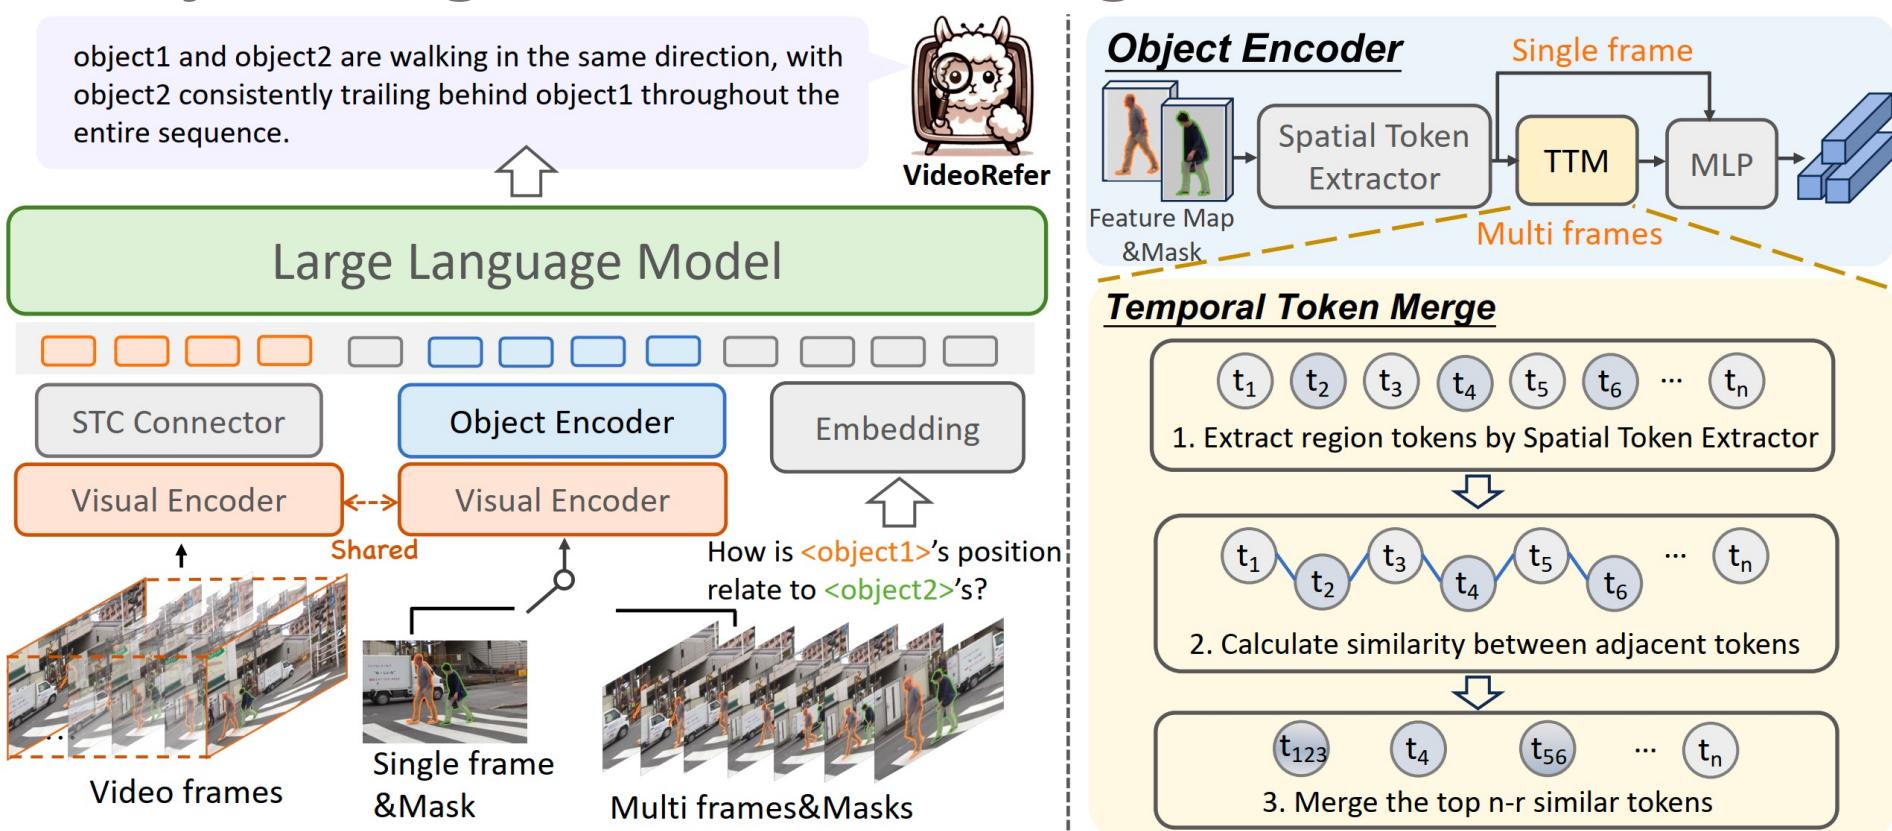
### VideoRefer-700K—Multi-agent Data Engine



### VideoRefer-Bench

# Fine-grained Object/Region Understanding

## VideoRefer Model



Base Model: VideoLLaMA2

A plug-and-play Spatial-Temporal Object Encoder:

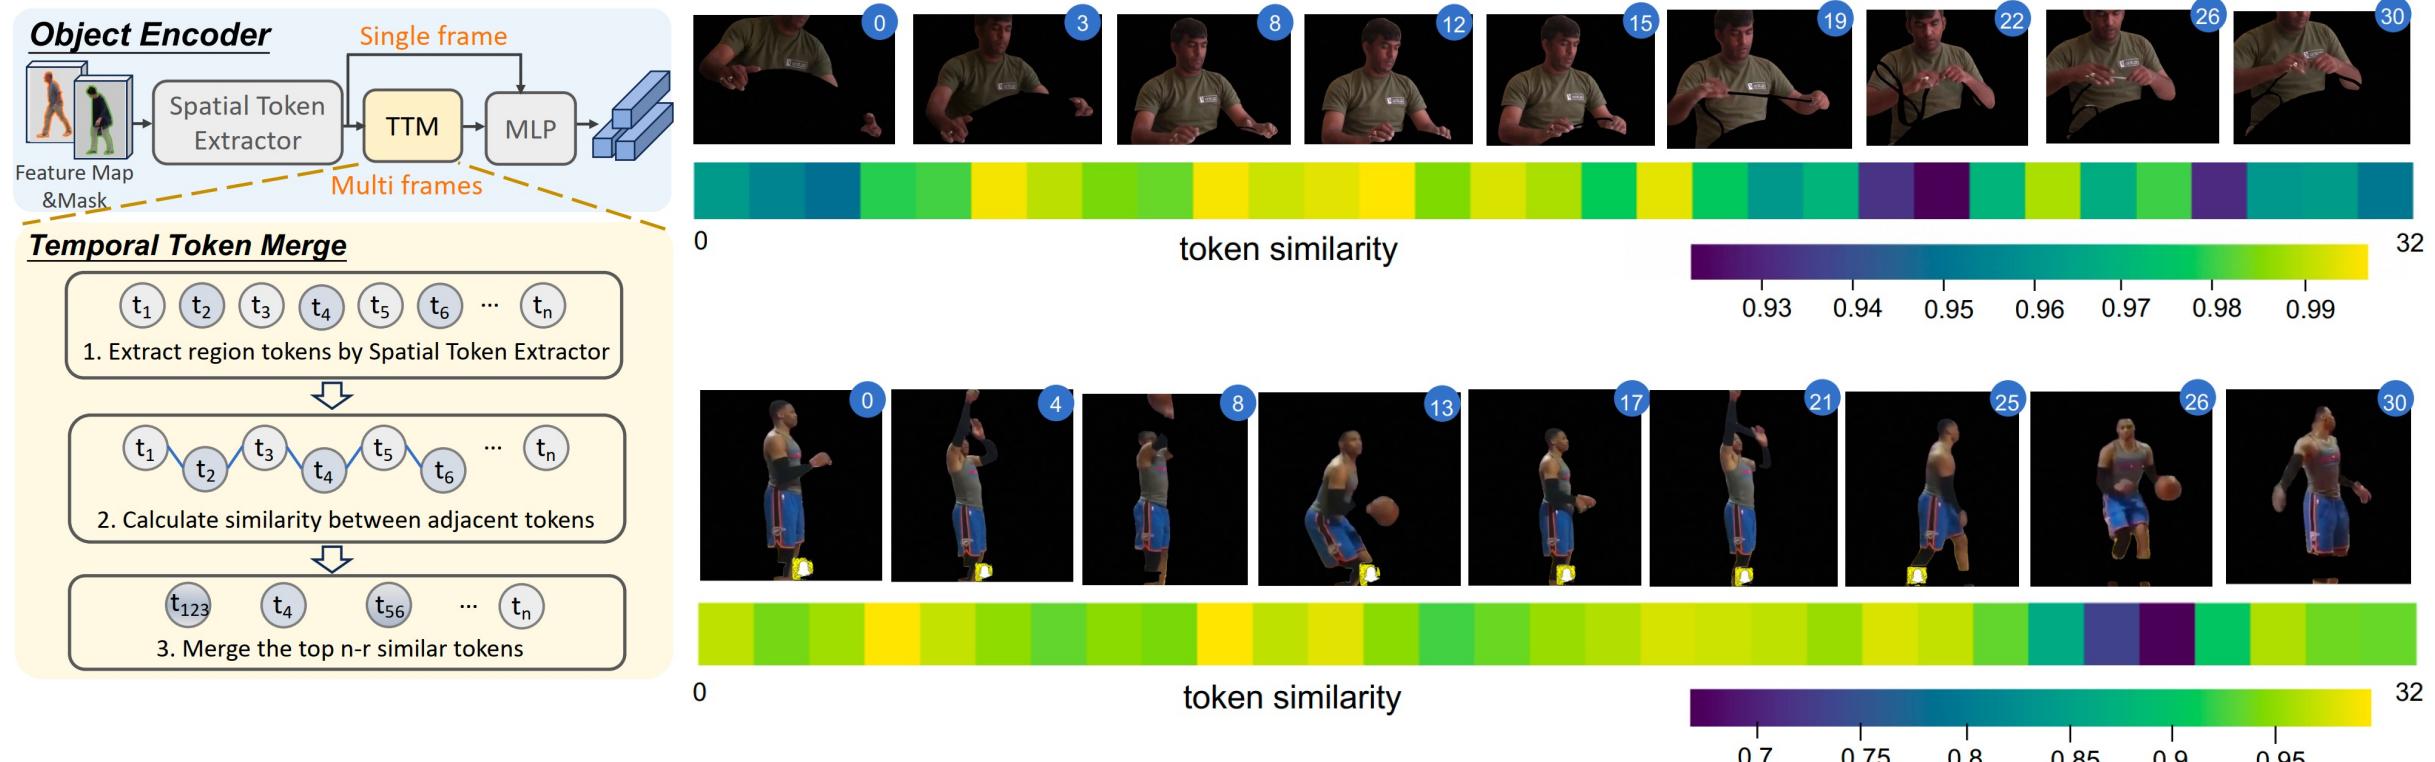
- Spatial Token Extractor (*Single-frame*)
- Temporal Token Merge Module (*Multi-frame*)
- Free-from input region (*Mask*)

Optimization Loss:

$$\mathcal{L} = \sum_{(V, R, x, y)} \log P(y | V, R_1, \dots, R_n, x)$$

# Fine-grained Object/Region Understanding

## VideoRefer Model

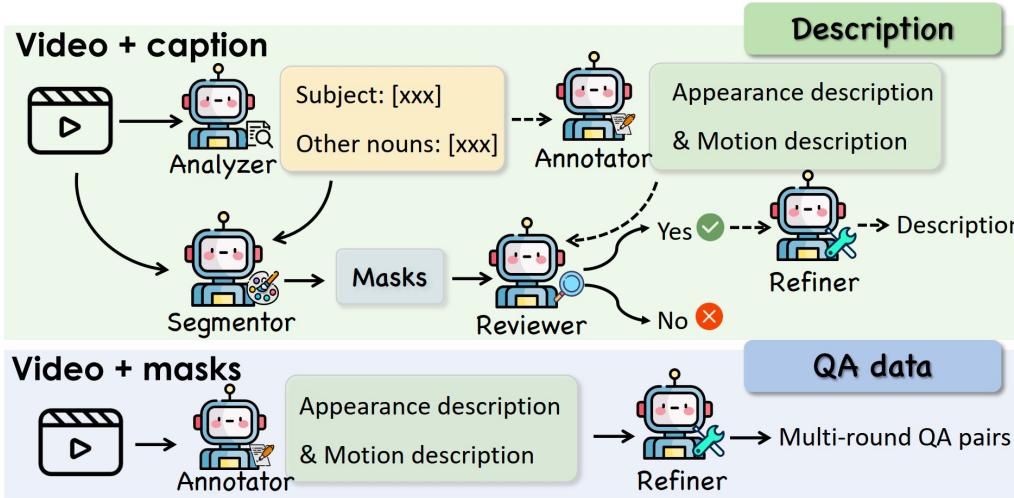


Compute the cosine similarity between each pair of adjacent tokens:

$$S_{m,m+1} = \frac{\mathbf{O}_m \cdot \mathbf{O}_{m+1}}{\|\mathbf{O}_m\| \cdot \|\mathbf{O}_{m+1}\|}, 0 \leq m < k$$

# Fine-grained Object/Region Understanding

## VideoRefer-700K



## Multi-agent Data Engine

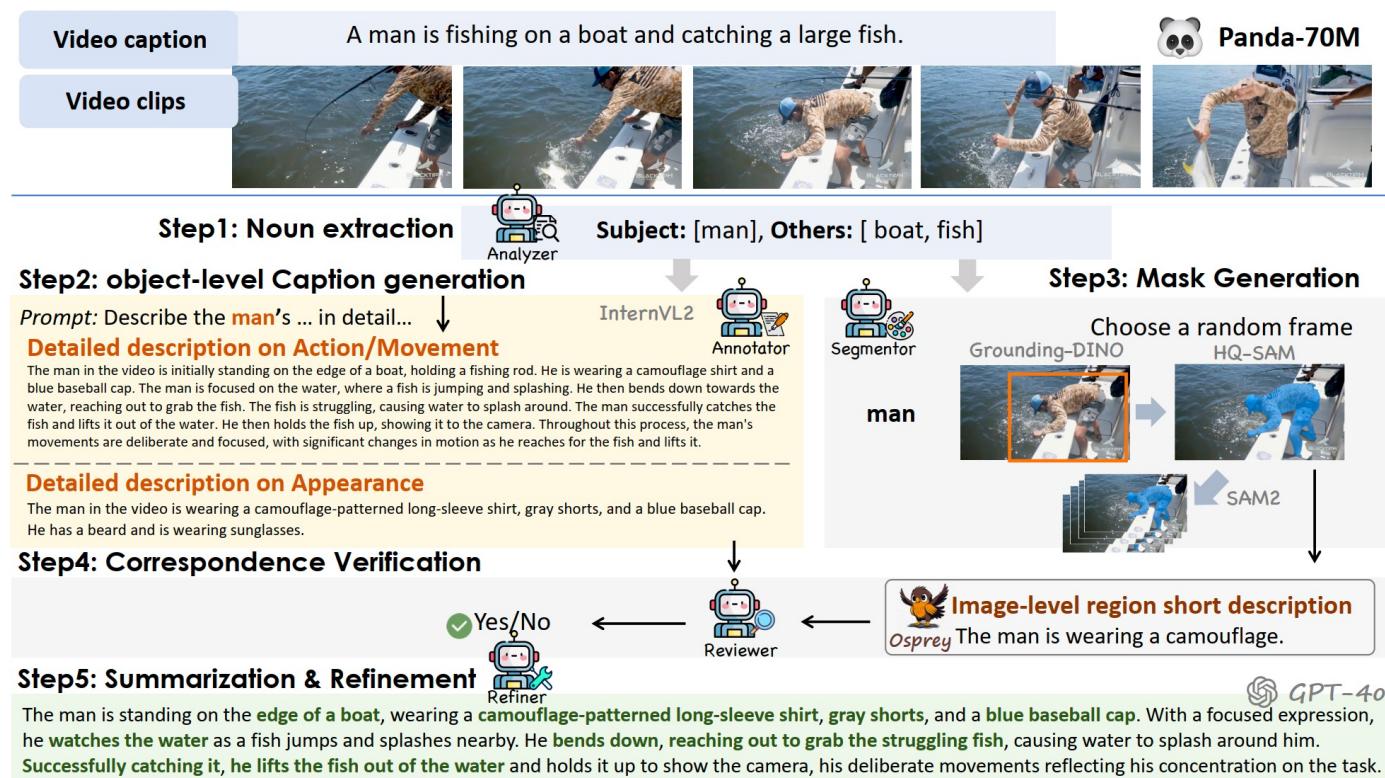
Step1- Analyzer: Qwen2-Instruct-7B

Step2-Annotator: InternVL2-26B

Step3-Segmentor:Grounding DINO&SAM 2

Step4-Reviewer: Osprey&Qwen2-Instruct-7B

Step5-Refiner:GPT-4o



## Three types:

- Object-level Detailed Caption
- Object-level Short Capton
- Object-level QA

	Manually True	Manually False
Reviewer True	88 (TP)	12 (FP)
Reviewer False	36 (FN)	64 (TN)

Table 8. Confusion matrix of the randomly sampled 100 items in the Reviewer evaluation.

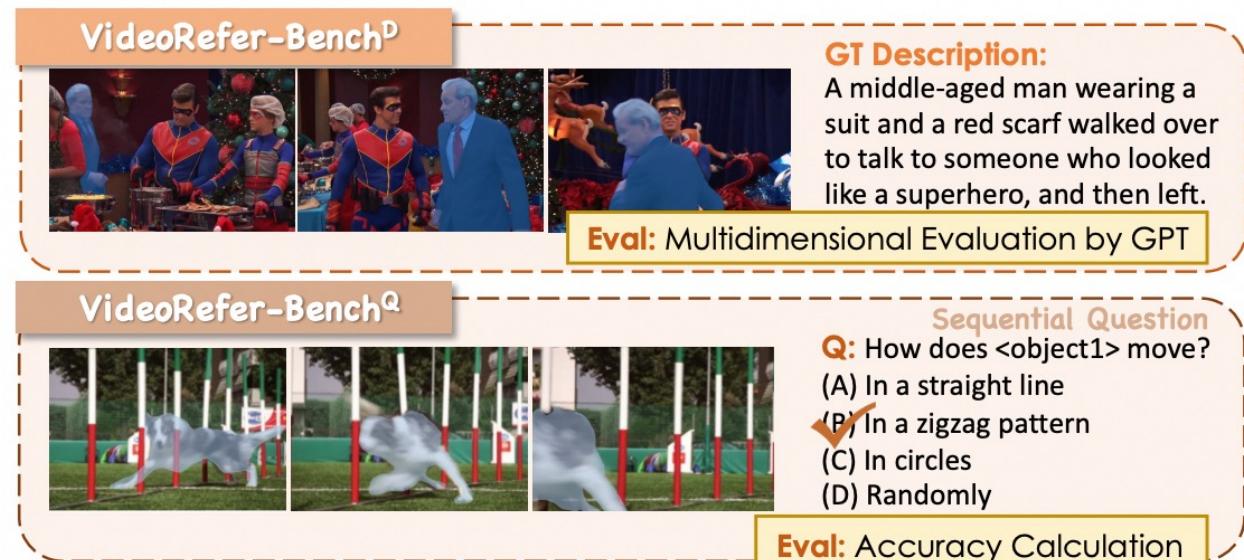
# Fine-grained Object/Region Understanding

## VideoRefer-Bench

### VideoRefer-Bench<sup>D</sup> (Description Generation)

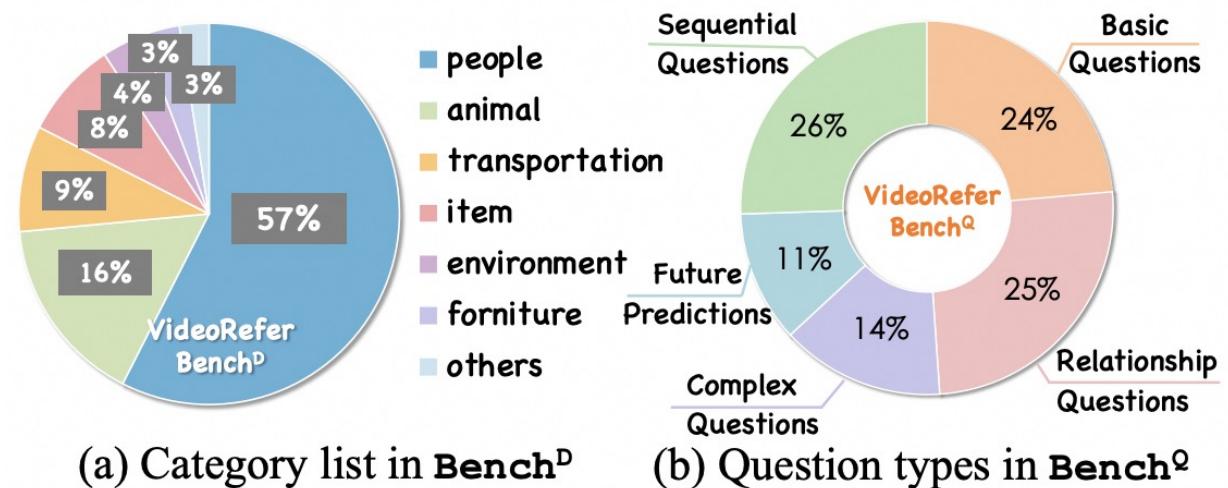
GPT assign scores from 0 to 5 across:

- Subject Correspondence
- Appearance Description
- Temporal Description
- Hallucination Detection



### VideoRefer-Bench<sup>Q</sup> (Multi-choice QA)

- Basic Questions
- Sequential Questions
- Relationship Questions
- Reasoning Questions
- Future Predictions



# Fine-grained Object/Region Understanding

## Experiments

Method	Single-Frame					Multi-Frame				
	SC	AD	TD	HD	Avg.	SC	AD	TD	HD	Avg.
<i>Generalist Models</i>										
LongVU-7B [38]	2.02	1.45	1.98	1.12	1.64	2.33	1.80	2.39	1.68	2.05
LongVA-7B [54]	2.63	1.59	2.12	2.10	2.11	3.02	2.30	1.92	2.51	2.44
LLaVA-OV-7B [15]	2.62	1.58	2.19	2.07	2.12	3.09	1.94	2.50	2.41	2.48
Qwen2-VL-7B [45]	2.97	2.24	2.03	2.31	2.39	3.30	2.54	2.22	2.12	2.55
InternVL2-26B [8]	3.55	2.99	2.57	2.25	2.84	4.08	3.35	3.08	2.28	3.20
GPT-4o-mini [29]	3.56	2.85	2.87	2.38	2.92	3.89	3.18	2.62	2.50	3.05
GPT-4o [29]	3.34	2.96	3.01	2.50	2.95	4.15	3.31	3.11	2.43	3.25
<i>Specialist Models</i>										
<i>Image-level models</i>										
Ferret-7B [46]	3.08	2.01	1.54	2.14	2.19	3.20	2.38	1.97	1.38	2.23
Osprey-7B [48]	3.19	2.16	1.54	2.45	2.34	3.30	2.66	2.10	1.58	2.41
<i>Video-level models</i>										
Elysium-7B [43]	2.35	0.30	0.02	3.59	1.57	—	—	—	—	—
Artemis-7B [33]	—	—	—	—	—	3.42	1.34	1.39	2.90	2.26
<b>VideoRefer-7B</b>	<b>4.41</b>	<b>3.27</b>	<b>3.03</b>	<b>2.97</b>	<b>3.42</b>	<b>4.44</b>	<b>3.27</b>	<b>3.10</b>	<b>3.04</b>	<b>3.46</b>

Method	Basic Questions	Sequential Questions	Relationship Questions	Reasoning Questions	Future Predictions	Average
<i>Generalist Models</i>						
LongVU-7B [38]	47.2	61.3	57.5	85.3	65.8	61.0
LongVA-7B [54]	56.2	62.5	52.0	83.9	65.8	61.8
InternVL2-26B [8]	58.5	63.5	53.4	88.0	78.9	65.0
GPT-4o-mini [29]	57.6	67.1	56.5	85.9	75.4	65.8
Qwen2-VL-7B [45]	62.0	69.6	54.9	87.3	74.6	66.0
LLaVA-OV-7B [15]	58.7	62.9	64.7	87.4	76.3	67.4
GPT-4o [29]	62.3	74.5	66.0	88.0	73.7	71.3
<i>Specialist Models</i>						
Osprey-7B [48]	45.9	47.1	30.0	48.6	23.7	39.9
Ferret-7B [46]	35.2	44.7	41.9	70.4	74.6	48.8
<b>VideoRefer-7B</b>	<b>75.4</b>	<b>68.6</b>	<b>59.3</b>	<b>89.4</b>	<b>78.1</b>	<b>71.9</b>

Method	Perception-Test	MVBench	VideoMME
VideoLLaMA2 [9]	51.4	54.6	47.9/50.3
VideoLLaMA2.1 [9]	54.9	57.3	54.9/56.4
Artemis [33]	47.1	34.1	28.8/35.3
<b>VideoRefer</b>	<b>56.3</b>	<b>59.6</b>	<b>55.9/57.6</b>

Mode	VideoRefer-Bench <sup>D</sup>			VideoRefer-Bench <sup>Q</sup>		
	TD	HD	Avg.	SQ	RQ	Avg.
Single-frame	3.03	2.97	3.42	68.3	59.1	71.9
Multi-frame	<b>3.10</b>	<b>3.04</b>	<b>3.46</b>	<b>70.6</b>	<b>60.5</b>	<b>72.1</b>

# Fine-grained Object/Region Understanding



## Describe Anything Model (DAM)

### Describe Anything: Detailed Localized Image and Video Captioning

Long Lian<sup>1,2</sup> Yifan Ding<sup>1</sup> Yunhao Ge<sup>1</sup> Sifei Liu<sup>1</sup> Hanzi Mao<sup>1</sup> Boyi Li<sup>1,2</sup> Marco Pavone<sup>1</sup>

Ming-Yu Liu<sup>1</sup> Trevor Darrell<sup>2</sup> Adam Yala<sup>2,3</sup> Yin Cui<sup>1</sup>

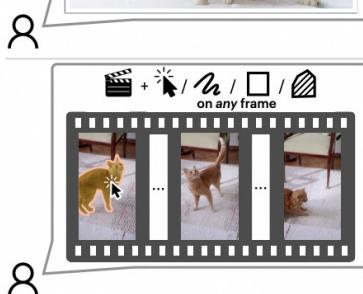
<sup>1</sup>NVIDIA

<sup>2</sup>UC Berkeley

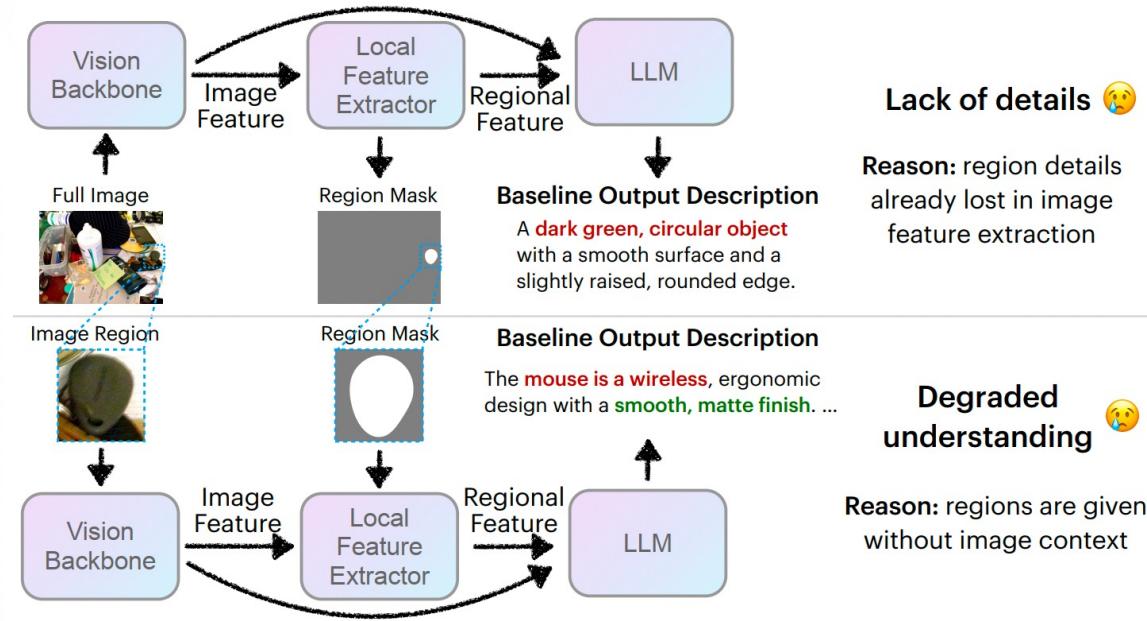
<sup>3</sup>UCSF



A white, fluffy dog with a thick coat, pointed ears, and a black nose. The dog has a wide-open mouth with a pink tongue hanging out, and its eyes are dark and alert. The fur is dense and appears soft, with some darker patches on its back.



A ginger cat with a sleek, short-haired coat is captured in a series of graceful movements. Initially, the cat is seen with its head turned slightly to the side, ears perked up, and eyes wide open, suggesting alertness or curiosity. As the sequence progresses, the cat's body is elongated, indicating a stretch or a poised stance. Its tail is held high, curving slightly at the tip, a sign of confidence or playfulness. The cat's front paw is extended forward, as if reaching or preparing to pounce.



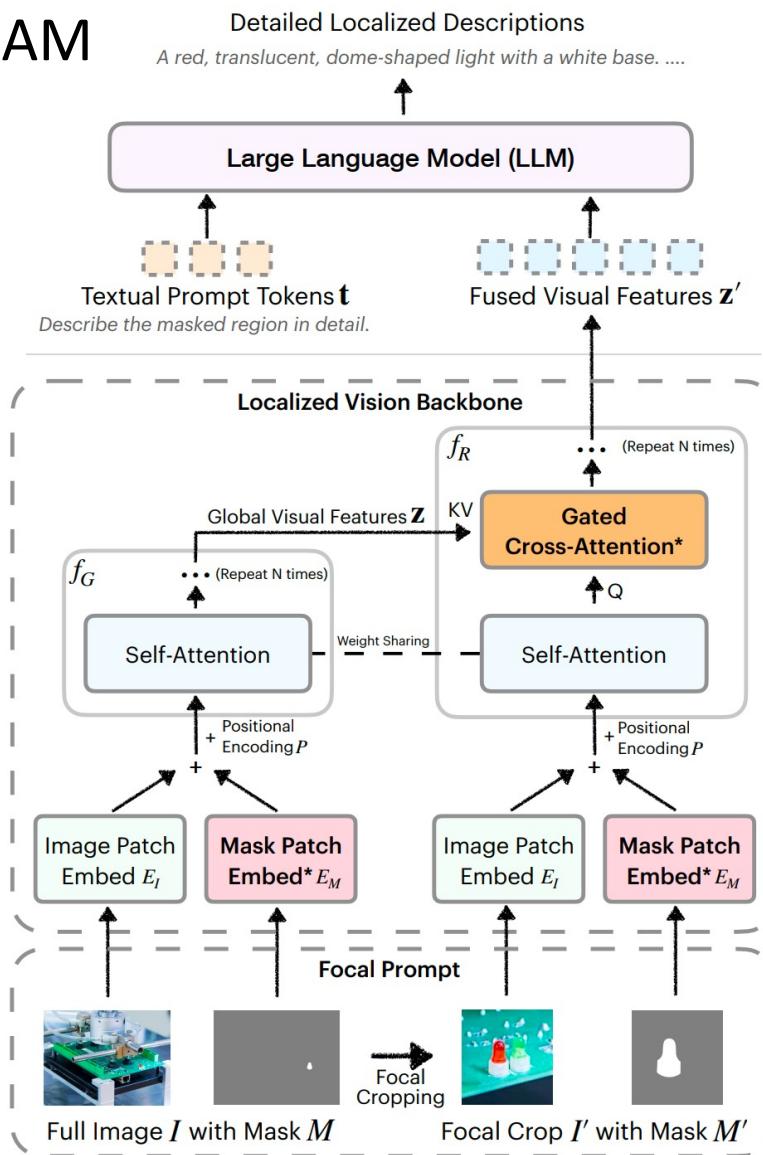
Typical two regional frameworks

Figure 1: Describe Anything Model (DAM) generates **detailed localized captions** for user-specified regions within **images** (top) and **videos** (bottom). DAM accepts various region specifications, including clicks, scribbles, boxes, and masks. For videos, specifying the region in *any frame* suffices.

Adopting VideoRefer-Bench & Osprey Evaluation.

# Fine-grained Object/Region Understanding

DAM



- Focal Prompt

**Full image and a zoomed-in region with corresponding mask**

$$x = E_I(I) + E_M(M) + P, \quad z = f_G(x)$$

$$x' = E_I(I') + E_M(M') + P, \quad z' = f_R(x', z)$$

- Localized Vision Backbone

Inject global features into the encoding of local regions using

## Gated Cross-Attention Adaptor

$$\mathbf{h}^{(l)'} = \mathbf{h}^{(l)} + \tanh(\gamma^{(l)}) \cdot \text{CrossAttn}(\mathbf{h}^{(l)}, \mathbf{z}),$$

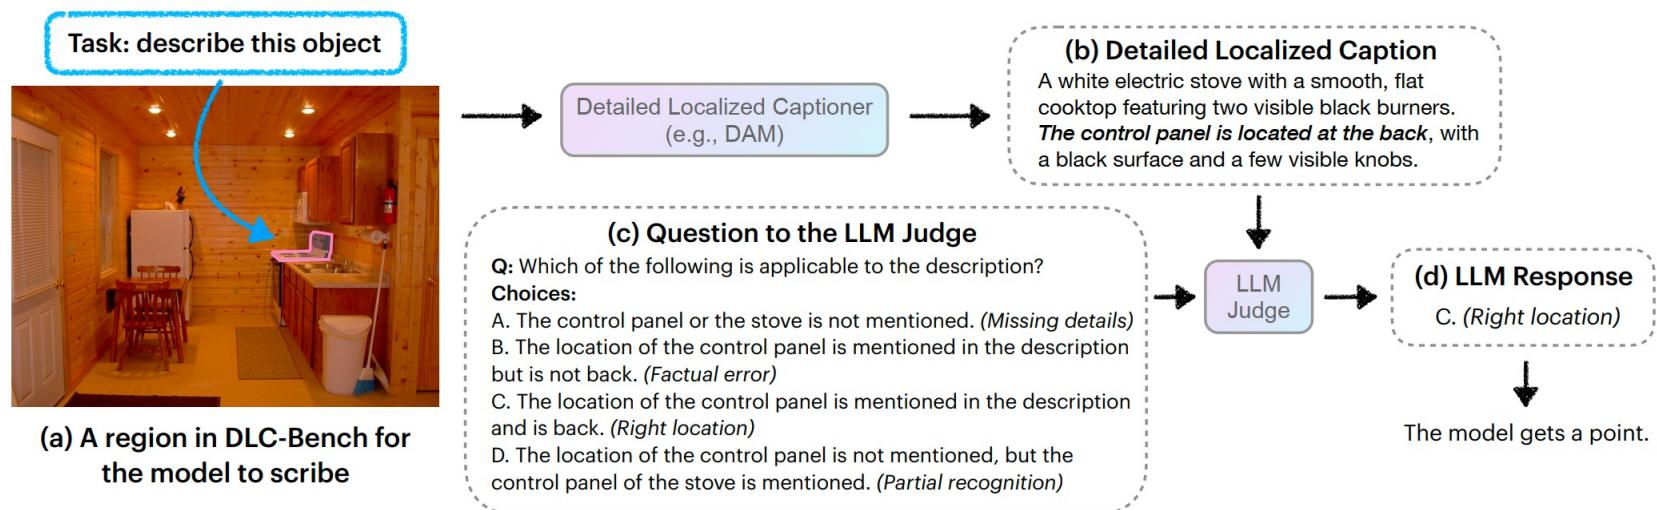
$$\mathbf{h}_{\text{Adapter}}^{(l)} = \mathbf{h}^{(l)'} + \tanh(\beta^{(l)}) \cdot \text{FFN}(\mathbf{h}^{(l)'}, \mathbf{z}),$$

# Fine-grained Object/Region Understanding

- **Simple Extension to Video Frames**

- All frames are naïvely concatenated along the temporal axis, **without considering inter-frame correlations**;
- Each object per frame is represented by **196 tokens**;
- Limited to captioning tasks only.

- Using LLM as Judge for Performance Evaluation & Dataset



Dataset	# Images	# Regions
<i>Stage 1:</i>		
LVIS [29]	90,613	373,551
Mapillary Vistas v2.0 [53]	17,762	100,538
COCO Stuff [11]	28,365	32,474
OpenImages v7 [33, 35]	64,874	96,006
PACO [60]	24,599	81,325
<i>Stage 2:</i>		
SA-1B (10%)	592,822	774,309
<b>Total</b>	<b>819,035</b>	<b>1,458,203</b>

# Fine-grained Object/Region Understanding

## Experiments

### The Evaluation setting of Osprey

Method	LVIS (%)		PACO (%)	
	Sem. Sim. (↑)	Sem. IoU (↑)	Sem. Sim. (↑)	Sem. IoU (↑)
	LLaVA-7B [48]	49.0	19.8	42.2
Shikra-7B [15]	49.7	19.8	43.6	11.4
GPT4RoI-7B [99]	51.3	12.0	48.0	12.1
Osprey-7B [95]	65.2	38.2	73.1	<u>52.7</u>
Ferret-13B [93]	65.0	37.8	-	-
VP-SPHINX-7B [45]	86.0	61.2	74.2	49.9
VP-LLAVA-8B [45]	<u>86.7</u>	<u>61.5</u>	<u>75.7</u>	50.0
<b>DAM-8B (Ours)</b>	<b>89.0</b>	<b>77.7</b>	<b>84.2</b>	<b>73.2</b>

Prompting	XAttn	#IT	Pos (%)	Neg (%)	Avg (%)
Full Image Only	No	196	32.1	65.4	48.7
Local Crop Only	No	196	43.5	76.6	60.1 ( <b>+11.4</b> )
Full + Local Crop	No*	392	26.3	58.6	42.4 ( <b>-6.3</b> )
Full + Local Crop	Yes	196	45.7	80.6	63.2 ( <b>+14.5</b> )
Focal Crop Only	No	196	47.3	<b>83.6</b>	65.4 ( <b>+16.7</b> )
<b>Full + Focal Crop</b>	Yes	196	<b>52.3</b>	82.2	<b>67.3 (+18.6)</b>

### VideoRefer-Bench

Method	SC	AD	TD	HD†	Avg.
<i>Zero-shot:</i>					
Qwen2-VL-7B [81]	3.30	2.54	2.22	2.12	2.55
InternVL2-26B [20]	4.08	<b>3.35</b>	3.08	2.28	3.20
GPT-4o-mini [54]	3.89	3.18	2.62	2.50	3.05
GPT-4o [54]	4.15	3.31	<b>3.11</b>	2.43	3.25
Osprey-7B [95]	3.30	2.66	2.10	1.58	2.41
Ferret-7B [93]	3.20	2.38	1.97	1.38	2.23
Elysium-7B [80]	2.35	0.30	0.02	<b>3.59</b>	1.57
Artemis-7B [59]	3.42	1.34	1.39	2.90	2.26
<b>DAM-8B (Ours)</b>	<b>4.45</b>	3.30	3.03	2.58	<b>3.34</b>
<i>In-domain*:</i>					
VideoRefer-7B [96]	4.44	3.27	3.10	3.04	3.46
<b>DAM-8B (Ours)</b>	<b>4.69</b>	<b>3.61</b>	<b>3.34</b>	<b>3.09</b>	<b>3.68</b>

Method	#Params	Pos (%)	Neg (%)	Avg (%)
<i>API-only General VLMs:</i>				
GPT-4o (SOM) [54]	-	5.0	29.2	17.1
o1 (SOM) [55]†	-	0.8	28.0	14.4
Claude 3.7 Sonnet (SOM) [73]†	-	0.5	40.2	20.4
Gemini 2.5 Pro (SOM) [74, 75]†	-	13.2	65.0	39.1
<i>Open-source General VLMs:</i>				
Llama-3.2 Vision (SOM) [25]	11B	16.8	40.4	28.6
Llama-3 VILA1.5 (SOM) [44]	8B	0.6	0.6	0.6
InternVL2.5 (SOM) [20, 21, 84]	8B	8.6	28.6	18.6
LLaVA v1.6 (SOM) [46–48]	7B	2.2	3.8	3.0
Qwen2.5-VL (SOM) [77, 81]	7B	8.5	27.2	17.8
VILA1.5 (SOM) [44]	3B	-0.4	15.4	7.5
<b>DAM (Ours)</b>	3B	<b>52.3</b>	<b>82.2</b>	<b>67.3</b>

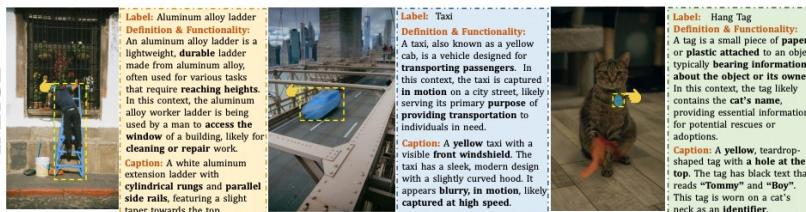
# Fine-grained Object/Region Understanding

## Perceive Anything: Recognize, Explain, Caption, and Segment Anything in Images and Videos

Weifeng Lin<sup>1\*</sup> Xinyu Wei<sup>3\*</sup> Ruichuan An<sup>4\*</sup> Tianhe Ren<sup>2\*</sup> Tingwei Chen<sup>1</sup>  
Renrui Zhang<sup>1</sup> Ziyu Guo<sup>1</sup> Wentao Zhang<sup>4</sup> Lei Zhang<sup>3</sup> Hongsheng Li<sup>1†</sup>

<sup>1</sup>CUHK <sup>2</sup>HKU <sup>3</sup>PolyU <sup>4</sup>Peking University

### IMAGE



### VIDEO

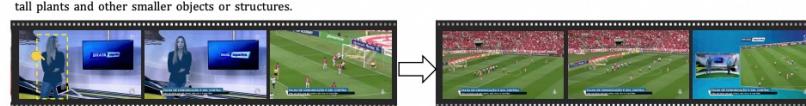
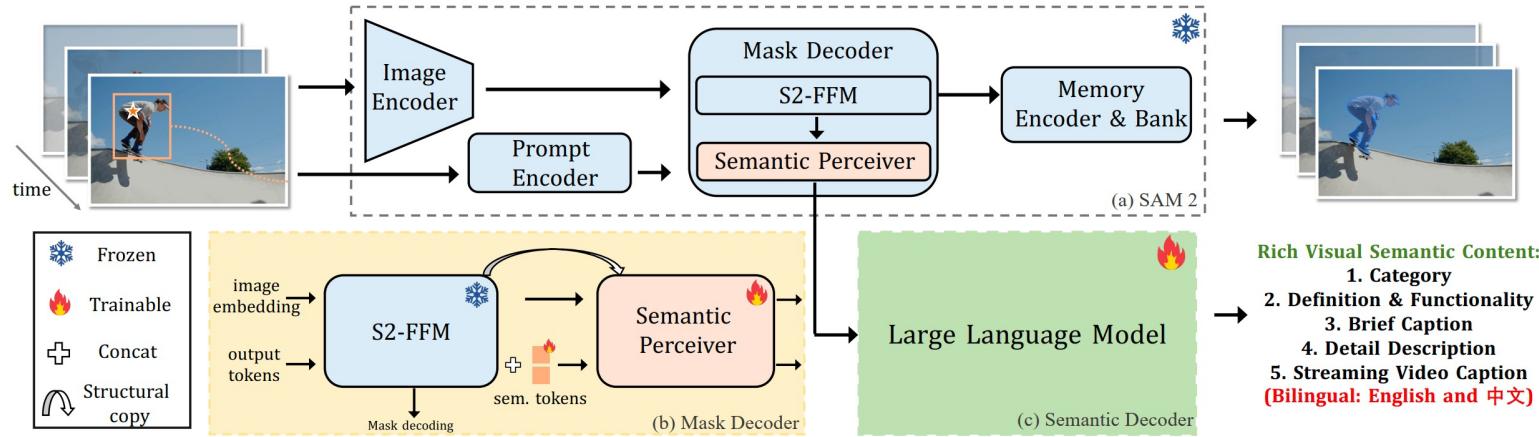


Figure 1: **Perceive Anything Model (PAM)**: PAM accepts various visual prompts (such as clicks, boxes, and masks) to produce region-specific information for images and videos, including masks, category, label definition, contextual function, and detailed captions. The model also handles demanding region-level streaming video captioning.

## Perceive Anything Model (PAM)

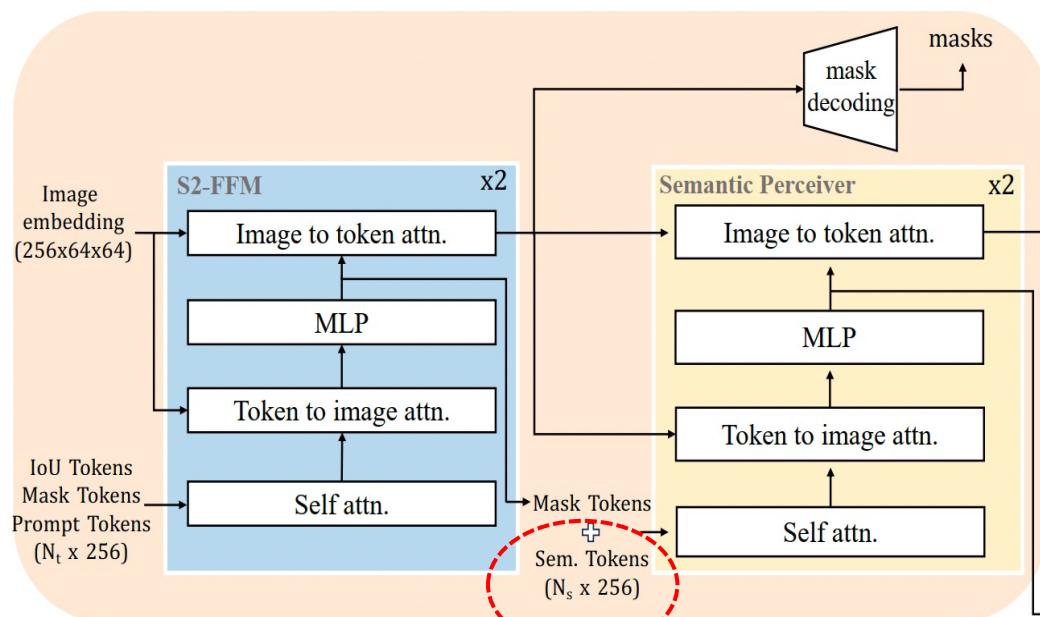


- Extends **SAM 2** by extracting its intermediate visual features and transforming them into **LLM-compatible tokens**.
- Enables segmentation mask decoding and semantic content decoding simultaneously.

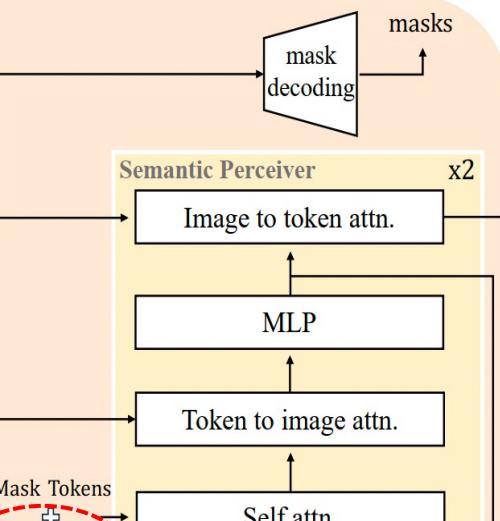
CUHK & HK PolyU

# Fine-grained Object/Region Understanding

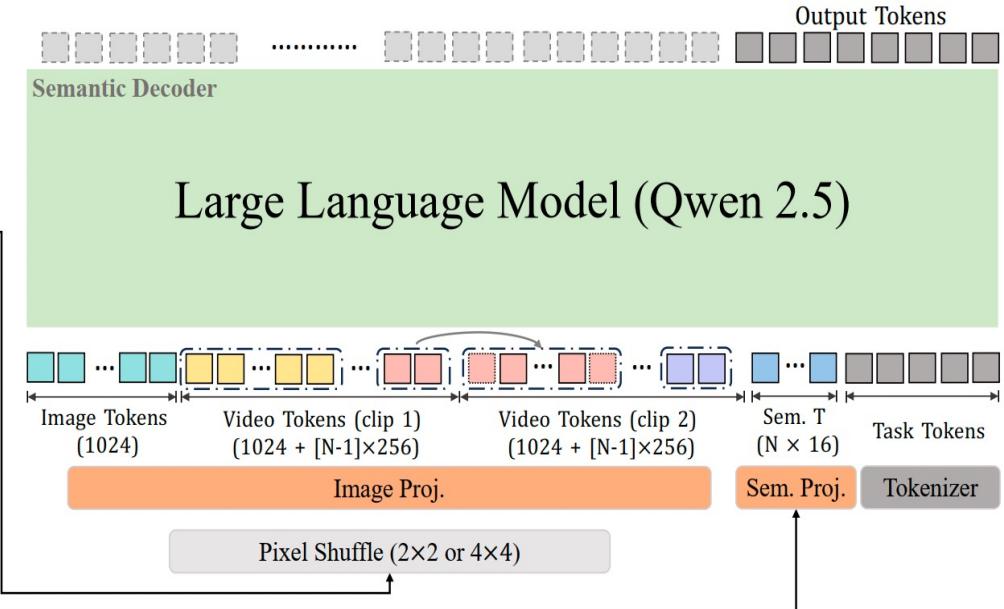
## S2-FFM



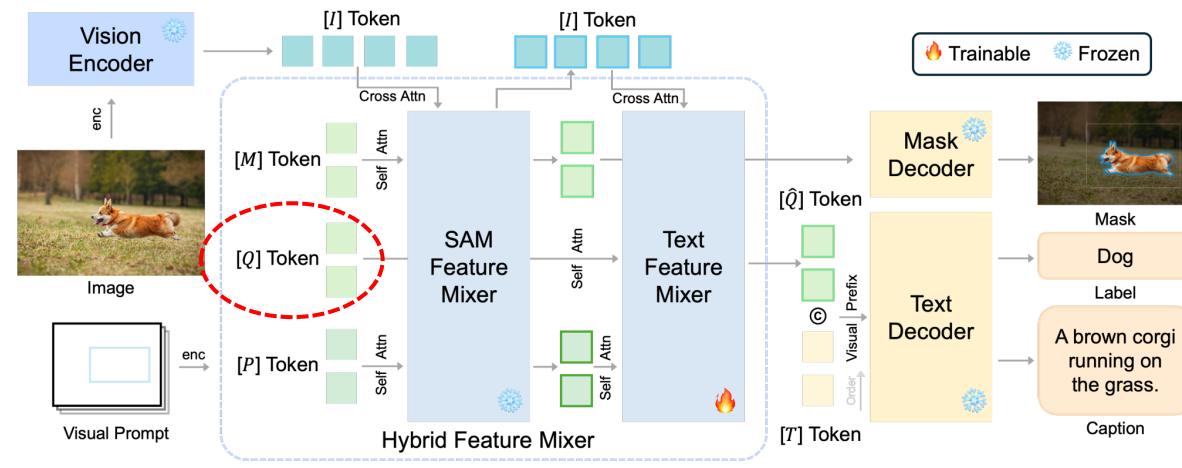
## Semantic Perceiver



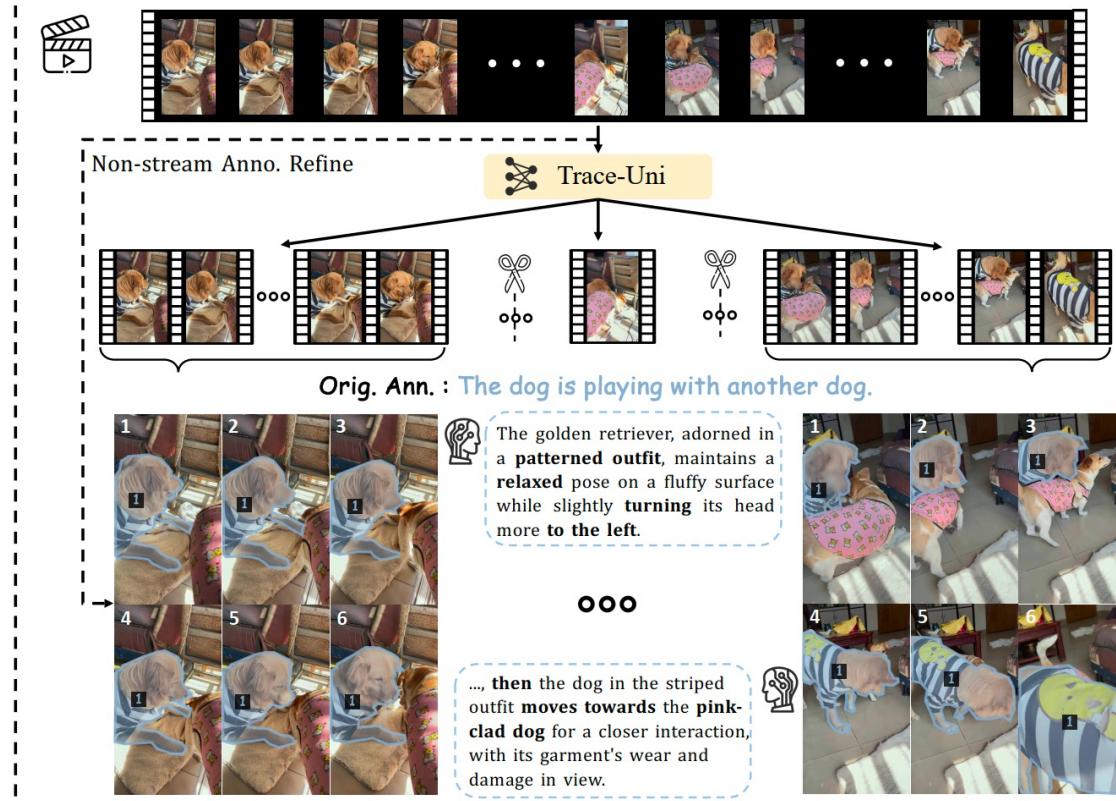
## 64x64 × N visual tokens & $N_s \times N$ semantic tokens



## SCA Model



# Fine-grained Object/Region Understanding



## Image Data:

1.5M image region-text pairs

## Video Data:

- Storyboard-based expansion
  - Event-aware segmentation
- 600K video region-text pairs

Supporting both English and Chinese.

A Large-Scale, Multi-Granular Region-Text Dataset

# Fine-grained Object/Region Understanding

## Streaming Object Caption



Limitations:

- **Fixed window** size without long-term memory;
- Limited to object captioning without multi-round, multi-object interaction.

# Content

1. Fine-grained Object/Region Understanding

Image/Video

**2. Efficient VLMs with Visual Token Compression**

**Model-driven:** TokenPacker, FastV, VisionZip, VisionTrim & LongVU

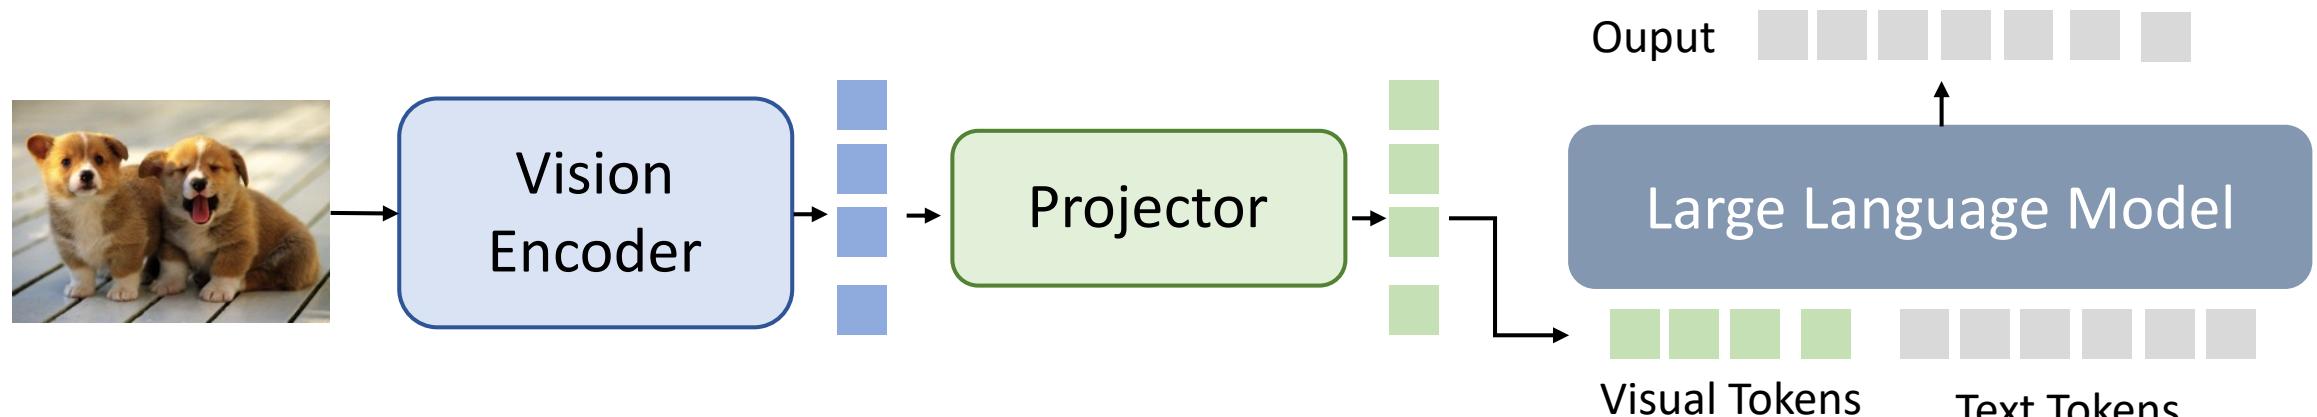
**Data-driven:** VocoLLAMA, Video-XL & DTR

**Other Paradigm:** mPLUG-Owl3, Lavi

3. Streaming Understanding & Interaction for AI Assistant

Training/Training-free

# Efficient VLMs with Visual Token Compression



- Vision Encoder
  - CLIP-VIT-L: ~0.3B
- Large Language Model
  - LLaMA/Vicuna: 7B/13B

- Visual Projector
  - MLP: 336x336 input -> 576 tokens



LLM dominates the main computational and memory demands.

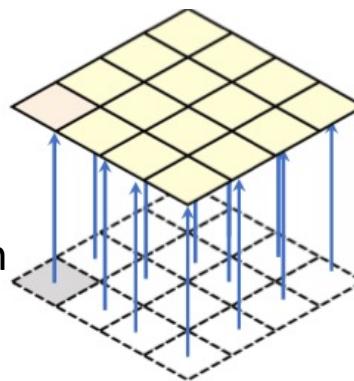


Reducing the number of visual tokens is a pivotal approach to bolster the efficiency.

# Efficient VLMs with Visual Token Compression

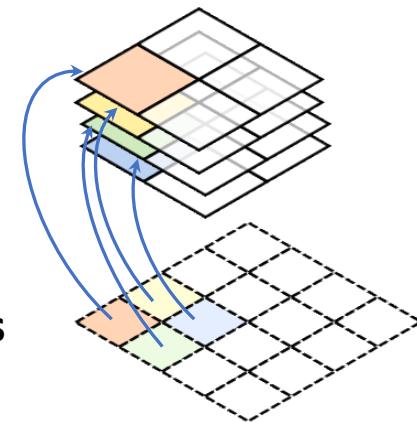
## Linear Projector

- One-to-one transformation
- 336x336 ->576token
- Retaining the detailed information with **redundant tokens**



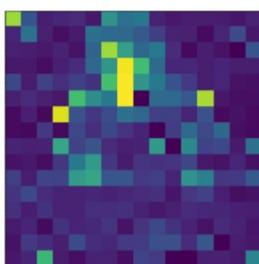
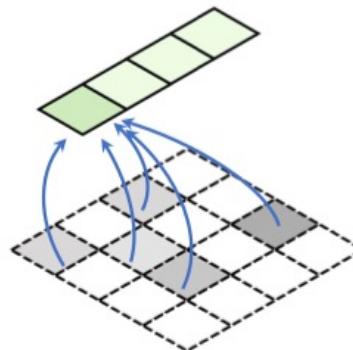
## Pixel shuffle

- Token reduction:144
- Nearby concatenation
- **Destroying intrinsic characteristics**



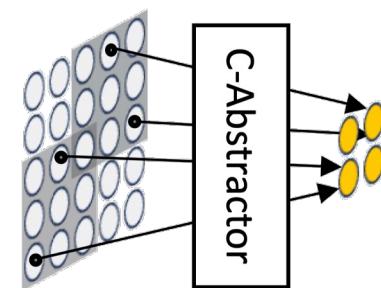
## Resampler/Q-Former

- Learnable queries (64/144)
- Extracting the most relevant visual tokens, **ignoring other objects.**



## Abstracter

- Local interaction
- Convolution layers
- **Omitting fine detailed information**



[1] Improved baselines with visual instruction tuning, in *NeurIPS2024*

[2] Qwen-vl: A frontier large vision-language model with versatile abilities, *Arxiv 2023*

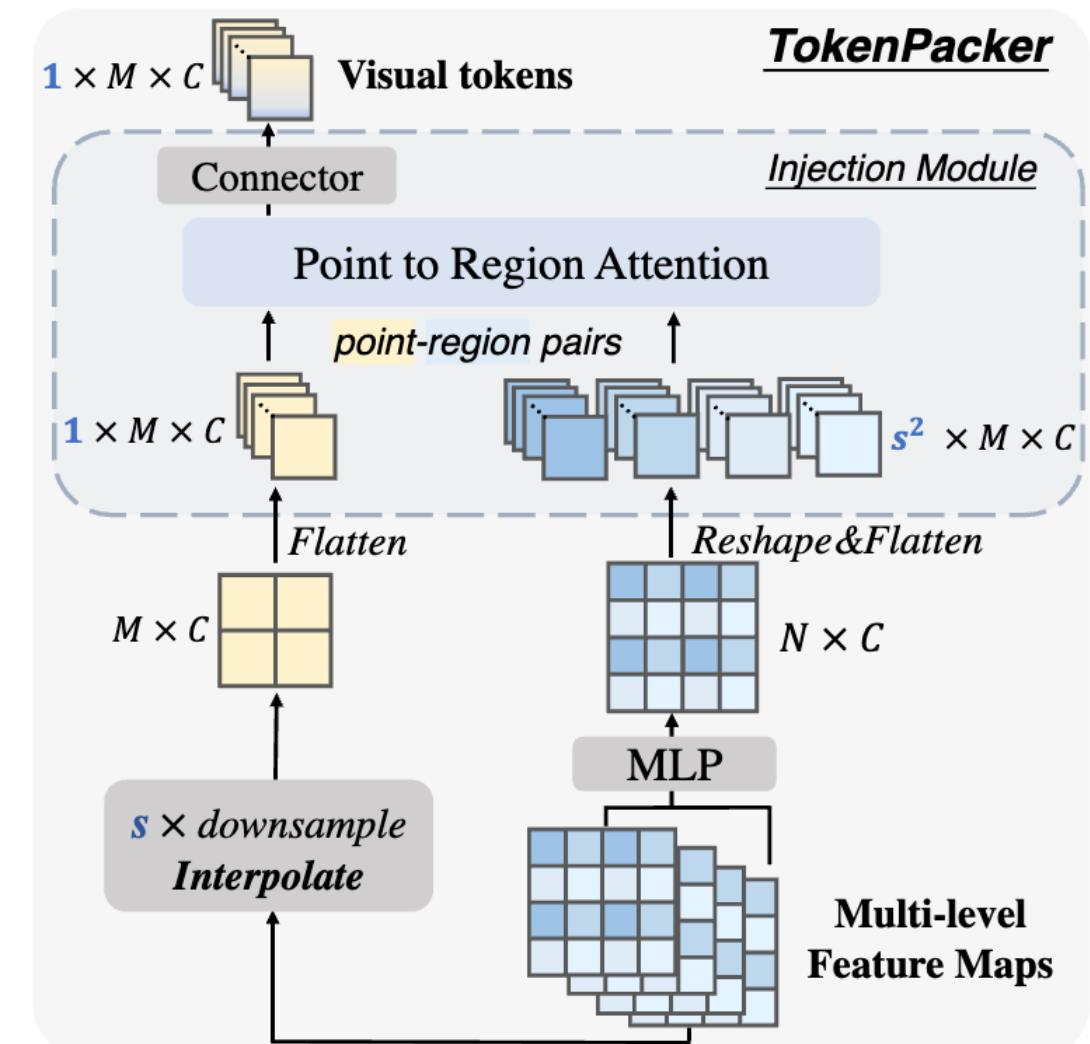
[3] How far are we to gpt-4v? closing the gap to commercial multimodal models with open-source suites, *Arxiv 2024*

[4] Honeybee: Locality-enhanced projector for multimodal llm, in *CVPR2024*

# Efficient VLMs with Visual Token Compression

## TokenPacker

- **Coarse-to-fine**  
Down-sampling features as coarse foundation
- **Point to Region Attention**, injecting the finer region feature to point query
- Multi-level visual features: 12-16-22-33
- **Scale factor**:  $S \in \{2,3,4\}$  to control the **reduction rate** {4, 9, 16}, even less.



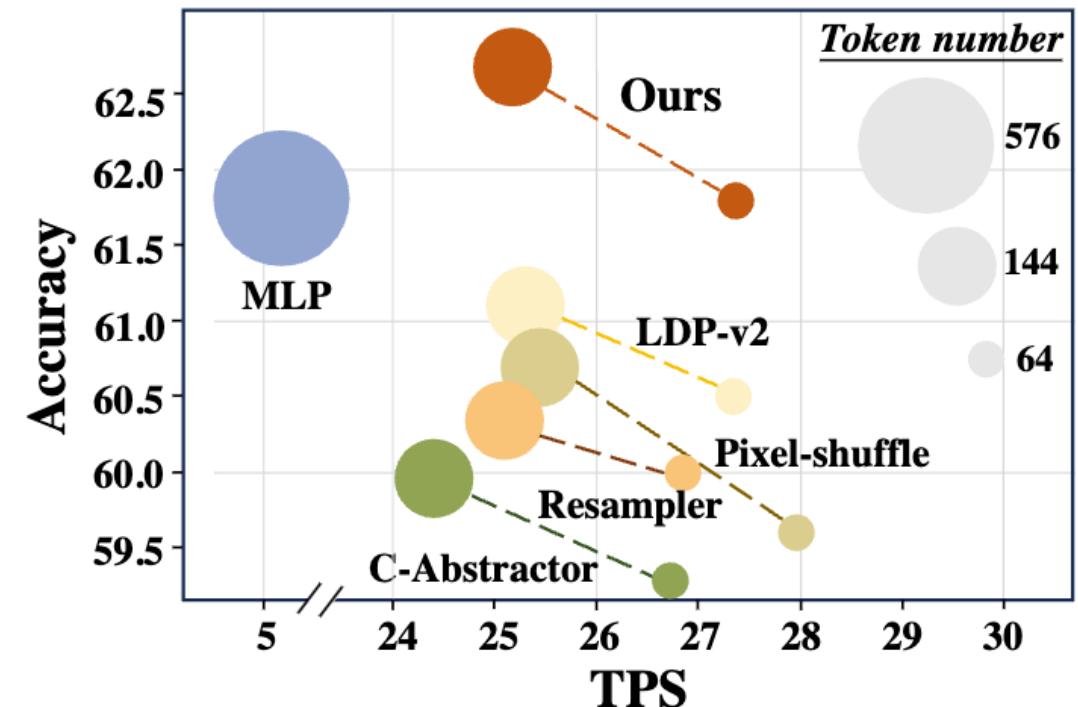
# Efficient VLMs with Visual Token Compression

Comparisons with same setting

LLaVA-1.5 as the baseline

Projector	#Token	TPS	MMB	MM-Vet	VQA <sup>v2</sup>	GQA	POPE	VizWiz	Avg.
MLP [9]	576	4.9	64.3	31.1	78.5	62.0	85.9	50.0	62.0
Resampler [11]	144	24.8	63.1	29.2	75.1	58.4	84.7	51.9	60.4
C-Abstractor [24]	144	24.1	63.1	29.4	74.6	59.2	84.6	49.2	60.0
Pixel-Shuffle [13]	144	<b>25.2</b>	64.0	29.7	76.2	60.1	85.9	48.8	60.8
LDP-v2 [26]	144	25.1	<b>66.2</b>	28.7	77.3	61.1	86.1	47.6	61.2
Ours	144	24.9	65.1	<b>33.0</b>	<b>77.9</b>	<b>61.8</b>	<b>87.0</b>	<b>52.0</b>	<b>62.8</b>
Resampler [11]	64	26.6	63.4	29.2	74.1	57.7	83.4	<b>53.0</b>	60.1
C-Abstractor [24]	64	26.5	62.5	29.0	74.4	59.3	62.5	45.6	59.3
Pixel-Shuffle [13]	64	<b>27.7</b>	63.2	28.5	74.6	59.1	85.2	47.4	59.7
LDP-v2 [26]	64	27.1	63.7	30.0	75.3	59.7	85.5	49.3	60.6
Ours	64	27.1	<b>64.1</b>	<b>31.7</b>	<b>77.2</b>	<b>61.1</b>	<b>86.3</b>	50.7	<b>61.9</b>

- TPS: token per second
- Evaluation on a NVIDIA A100 GPU

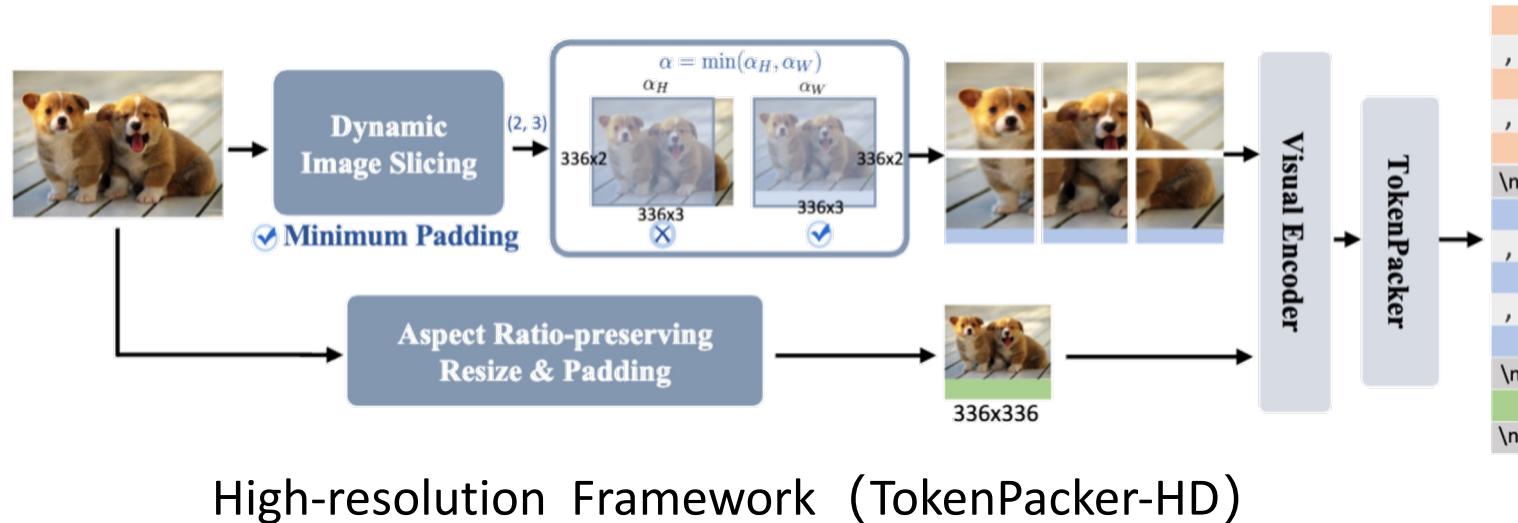


1/9 of the original results in a 5.5× acceleration, while maintaining comparable performance.

Exhibit a more favorable superiority on accuracy and efficient against other counterparts.

# Efficient VLMs with Visual Token Compression

## TokenPacker-HD



High-resolution Framework (TokenPacker-HD)

Method	LLM	#Data	Max Res.	#Token	VQA <sup>T</sup>	OCRB	DocVQA	MMB	MMMU	MME	VQA <sup>v2</sup>	VizWiz	POPE
OtterHD [26]	Fuyu-8B [5]	-	1024×1024	-	-	-	-	58.3	-	1294/-	-	-	86.0
SPHINX-2k [32]	LLaMA-13B	1.0B	762×762	2890	61.2	-	-	65.9	-	1471/-	80.7	44.9	87.2
UReader [56]	LLaMA-13B	86M	896×1120	-	57.6	-	<b>65.4</b>	-	-	-	-	-	-
Monkey [30]	QWen-7B	1.0B	896×1344	1792	-	<b>514</b>	-	-	-	-	80.3	<b>61.2</b>	67.6
TextHawk [59]	InternLM-7B	115M	1344×1344	-	-	-	<b>76.4</b>	<b>74.6</b>	-	1500/-	-	-	-
LLaVA-UHD [55]	Vicuna-13B	1.2M	672×1008	-	67.7	-	-	68.0	-	1535/-	81.7	56.1	<b>89.1</b>
LLaVA-NeXT [34]	Vicuna-7B	1.3M	672×672	2880	64.9	-	-	67.4	35.8	1519/332	81.8	57.6	86.5
LLaVA-NeXT [34]	Vicuna-13B	1.3M	672×672	2880	67.1	-	-	<b>70.0</b>	36.2	1575/326	<b>82.8</b>	60.5	86.2
Mini-Gemini-HD [28]	Vicuna-7B	2.7M	1536×1536	2880	68.4	456*	65.0*	65.8	36.8	1546/319	80.3*	54.6*	86.8*
Mini-Gemini-HD [28]	Vicuna-13B	2.7M	1536×1536	2880	<b>70.2</b>	501*	<b>70.0*</b>	68.6	37.3	1575/326	81.5*	57.2*	87.0*
LLaVA-TokenPacker-HD	Vicuna-7B	2.7M	1088×1088	~954 <sup>†</sup>	68.0	452	60.2	67.4	35.4	1489/338	81.2	54.7	88.2
LLaVA-TokenPacker-HD	Vicuna-13B	2.7M	1088×1088	~954 <sup>†</sup>	69.3	498	63.0	69.5	<b>38.8</b>	<b>1595/356</b>	<b>82.0</b>	59.2	88.1
LLaVA-TokenPacker-HD	Vicuna-13B	2.7M	1344×1344	~1393 <sup>†</sup>	<b>70.6</b>	<b>521</b>	<b>70.0</b>	68.7	37.4	1574/350	81.7	57.0	88.0
LLaVA-TokenPacker-HD	Vicuna-13B	2.7M	1344×1344	~619 <sup>#</sup>	68.8	470	63.0	69.9	<b>38.2</b>	<b>1577/353</b>	81.7	<b>61.0</b>	87.6
LLaVA-TokenPacker-HD	Vicuna-13B	2.7M	1344×1344	~347 <sup>\$</sup>	68.4	447	58.0	68.3	36.9	<b>1577/332</b>	81.2	58.1	88.0

Adopt the same training data Mini-Gemini[1]

Method	Res.	LVIS		PACO		Cityscapes	ADE20K
		SS	S-IoU	SS	S-IoU	AP	AP
Osprey [60]	512	65.2	38.2	73.1	52.7	29.2	31.8
FixedSplit [36]	672	69.4	45.6	79.3	63.5	33.7	39.5
AdaptiveSplit-1 [37]	Any	69.7	45.9	79.3	63.9	38.0	40.6
AdaptiveSplit-2 [56]	Any	70.0	46.3	79.3	63.9	42.3	41.0
Ours	Any	<b>71.6</b>	<b>47.5</b>	<b>79.8</b>	<b>64.1</b>	<b>43.8</b>	<b>42.0</b>

Employing Osprey on TokenPacker-HD framework

# Efficient VLMs with Visual Token Compression

FastV

(Within LLM Decoding)

Large Language Model  $M$

Output Tokens

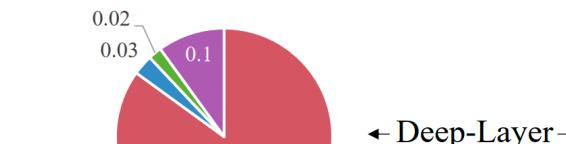
Training-free

System Prompt      Image      Instruction

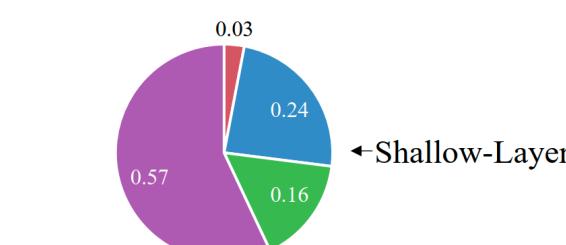
Attention Allocation:

$$\lambda_{sys}^j = \sum_{i=1}^n \alpha_{sys}^{i,j}$$

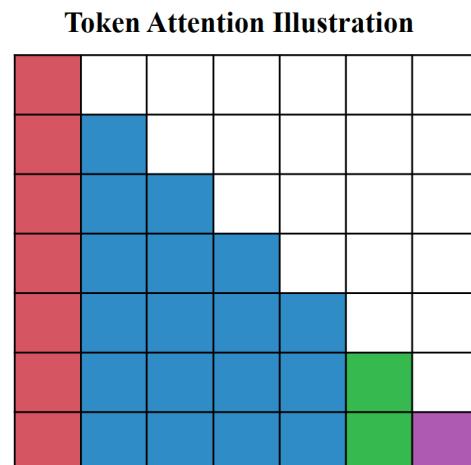
- System Prompt : 472x
- Image Tokens : 1x
- User Instruction: 3x
- Output Tokens : 12.8x



← Deep-Layer →



← Shallow-Layer →



■ System Prompt (35)   ■ User Instruction (135)  
■ Image Tokens (576)   ■ Output Tokens (150)

Attention Allocation  
(Decoding Output Tokens)

Attention Efficiency  
(Attention Allocation / Token Number)

# Efficient VLMs with Visual Token Compression

Training-free

FastV

(Within LLM Decoding)

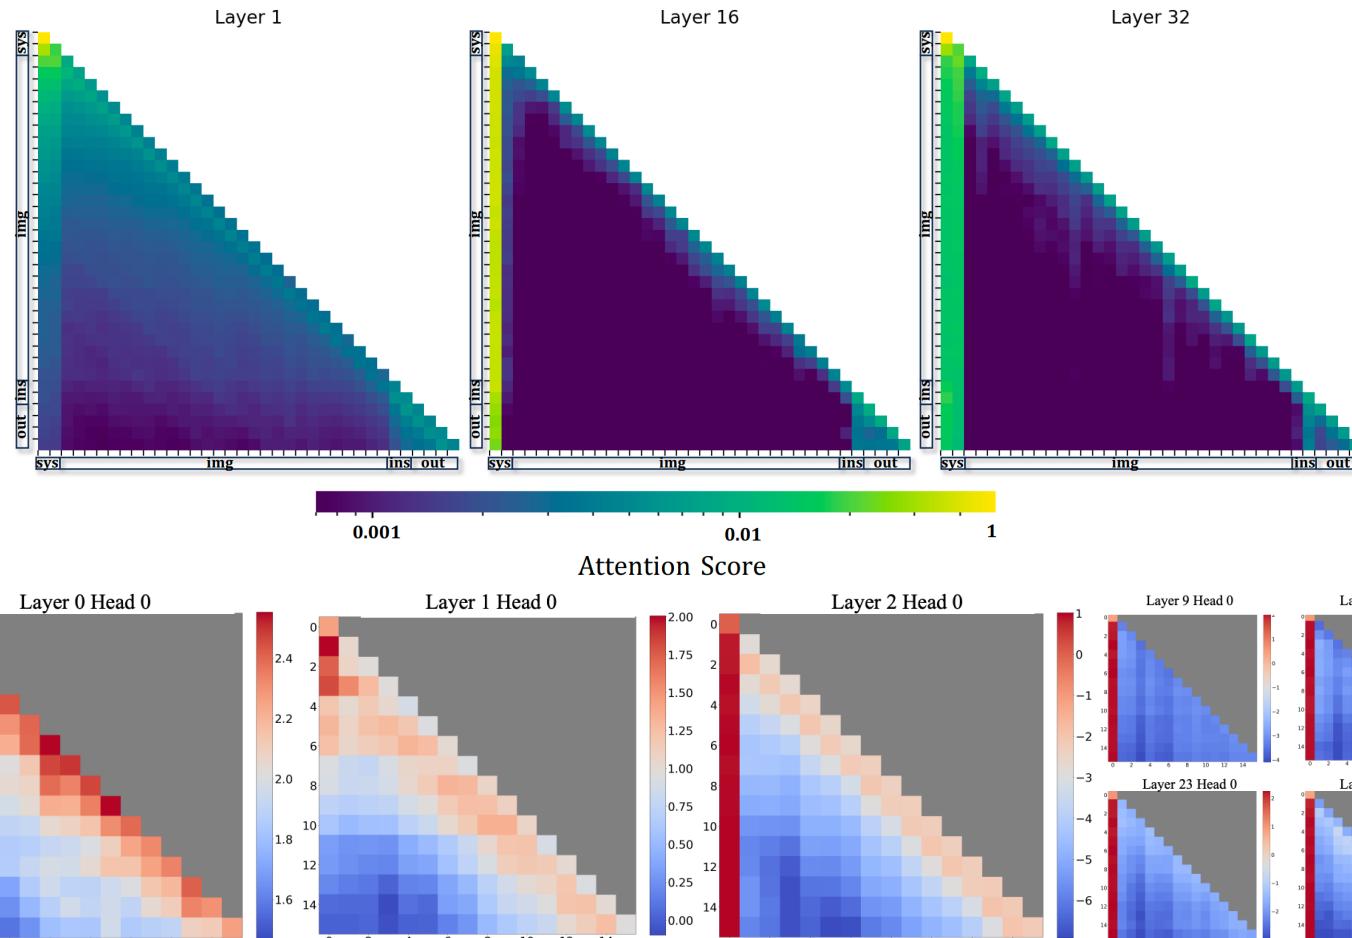


Figure 2: Visualization of the *average* attention logits in Llama-2-7B over 256 sentences, each with a length of 16. Observations include: (1) The attention maps in the first two layers (layers 0 and 1) exhibit the "local" pattern, with recent tokens receiving more attention. (2) Beyond the bottom two layers, the model heavily attends to the initial token across all layers and heads.

# Efficient VLMs with Visual Token Compression

Training-free

FastV

(Within  
LLM Decoding)



What is funny about the video?

User

The video is funny because the baby is trying to read a book while sitting on the bed, and he is pointing to different parts of the book. The baby's actions are amusing because he is not able to read the book properly, and he is pointing to random parts of the book.

The video is funny because the baby is trying to read a book while sitting on the bed, and he is pointing to different parts of the book. The baby's actions are amusing because he is not able to read the book properly, and he is pointing to random parts of the book.

The video is funny because the baby is trying to read a book while sitting on the bed, and he is pointing to different parts of the book. The baby's actions are amusing because he is not able to read the book properly, and he is pointing to random parts of the book.

The video is funny because the baby is pretending to read a book while making funny faces and gestures.

Identical

No FastV  
100% FLOPs

FastV(K=2, R=50%)  
52% FLOPs

FastV(K=5, R=75%)  
38% FLOPs

FastV(K=2, R=75%)  
33% FLOPs

FastV

Transformer Block K+2

Transformer Block K+1

FastV Re-rank & Filtering R%

Transformer Block K

8 \* 256 = 2048 Image Tokens



VideoLLaVA  
+ FastV

Image Tokens Text Tokens  
Filtered Image Tokens

CLIP-VIT + Projector

# Efficient VLMs with Visual Token Compression

Training-free

## Comparisons with same setting

Model	FastV Settings		Nocaps CIDEr	Flickr30k CIDEr	A-OKVQA Accuracy	MMMU Accuracy	Avg
	K	R					
LLaVA-1.5-7B	Baseline	99.3	100%	99.8	67.9	76.7	34.8 <b>69.8</b>
	2 90%	19.9	20%	72.1	43.7	70.1	35 55.2
	2 75%	32.8	33%	94.6	63.6	75.5	34.8 67.1
	2 50%	54.6	55%	99.7	67.5	77	34.4 <b>69.7</b>
	3 90%	22.8	23%	87.2	55.8	71.9	34.8 62.4
	3 75%	34.8	35%	98	65	74.7	34.1 68.0
	3 50%	56.6	57%	99.7	68.3	76.7	34.3 <b>69.8</b>
	5 90%	27.8	28%	88.6	59.3	70.6	33.9 63.1
	5 75%	39.7	40%	98.5	66.3	74.8	34.3 68.5
	5 50%	59.6	60%	99.2	67.9	76.8	34.3 69.6
	0 90%	18.9	19%	7	53.2	66.8	34.7 40.4
	0 75%	28.8	29%	27.2	61.4	72.8	35.1 49.1
	0 50%	51.6	52%	100.9	65.5	75.3	34.3 69.0
LLaVA-1.5-13B	Baseline	154.6	100%	102.8	73	82	36.4 <b>73.6</b>
	2 90%	29.7	19%	87.9	62	75	36.3 65.3
	2 75%	50.2	32%	100.5	72.5	80.9	38.1 73.0
	2 50%	84.6	55%	103.1	73.4	81	36.7 <b>73.6</b>
	3 90%	33.0	21%	90.2	63.6	75.2	34.9 66.0
	3 75%	52.9	34%	100.9	72.1	79.5	36.4 72.2
	3 50%	86.4	56%	102.7	73.4	81.3	36.4 <b>73.5</b>
	5 90%	39.6	26%	93.5	67.4	75.8	35.4 68.0
	5 75%	58.4	38%	101.4	72.5	80	36.2 72.5
	5 50%	90.1	58%	102.5	73.5	81.2	36.6 <b>73.5</b>
QwenVL-Chat-7B	Baseline	71.9	100%	94.9	72.5	75.6	35.8 <b>69.7</b>
	2 90%	15.8	22%	81.9	61.5	68.5	35.3 61.7
	2 75%	24.4	34%	90.5	67.0	75.1	35.3 67.0
	2 50%	39.5	55%	94.4	71.4	75.3	35.6 <b>69.2</b>



What are these? Describe the image in details



LLaVA1.5



LLaVA1.5+FastV

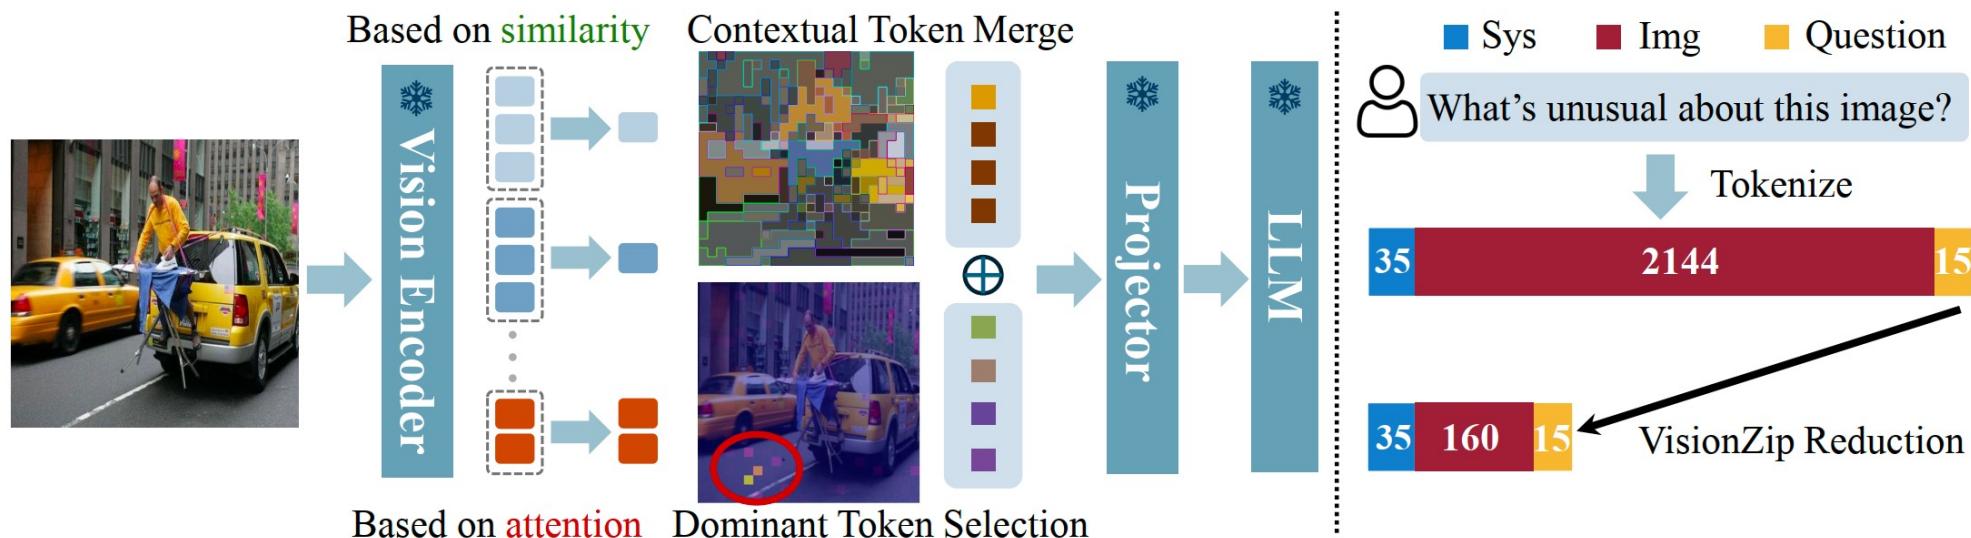
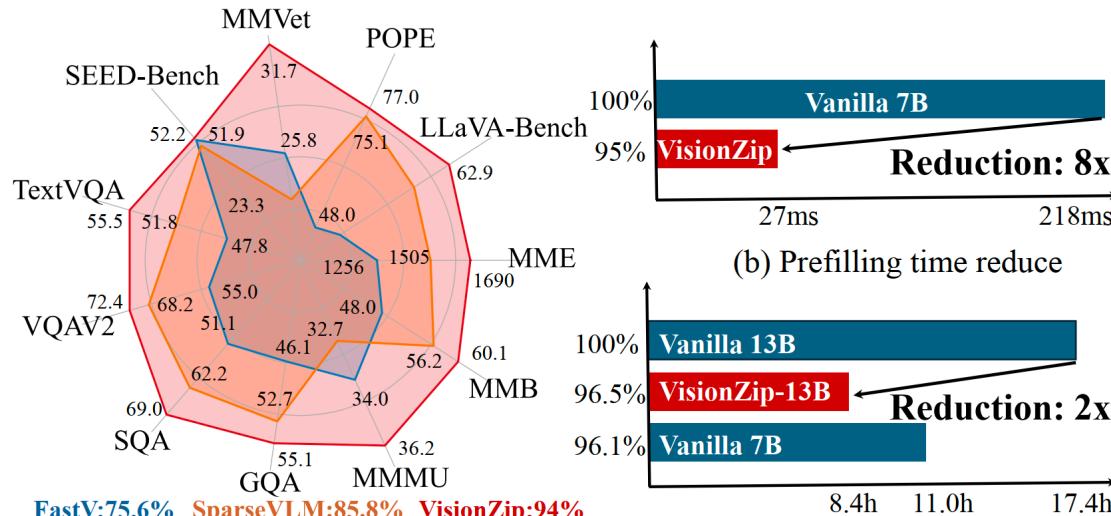


# Efficient VLMs with Visual Token Compression

Training-free

VisionZip

(Within Visual Encoding)

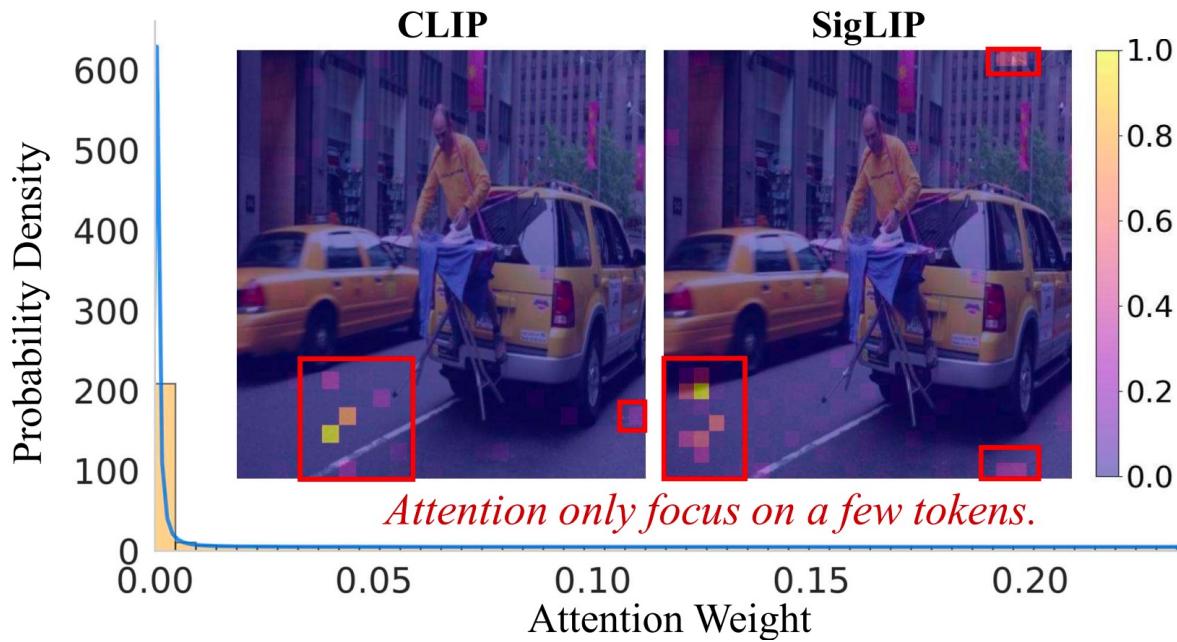


# Efficient VLMs with Visual Token Compression

Training-free

## Dominant Token Selection

- Using **[CLS] Tokens** attention scores to identify key visual tokens (CLIP)
- Average attention each token receives from all others (SigLIP)



### Algorithm 1 Pseudocode for Dominant Token Selection

```
# B: batch size; S: sequence length
# H: number of attention heads;
# K: number of target dominant tokens
# CLS_IDX: Index of the CLS token
# SELECT_LAYER: Selected layer for Visual Token

# set the output_attentions=True to get the attention
output = vision_tower(images, output_hidden_states=
    True, output_attentions=True)

#attn in shape (B, H, S, S)
attn = output.attentions[SELECT_LAYER]

#attn in shape (B, H, S, S)
vanilla_tokens = output.hidden_states[SELECT_LAYER]

#The attention received by each token
#If no CLS, use mean calculate received attention
attn_rec = attn[:, :, cls_idx, cls_idx+1:].sum(dim=1)

# Select K Dominant Tokens
_, topk_idx = attn_rec.topk(K, dim=1)

# Concat with the CLS token
dominant_idx = cat(CLSS_ID, topk_idx+1)

# filter the Dominant Tokens
dominant_tokens = vanilla_tokens.filter(dominant_idx)

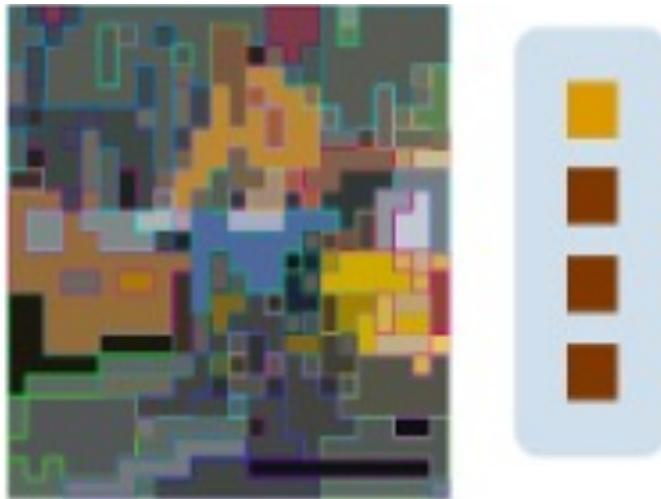
cat: concatenation; filter: select the tokens based on the index.
```

# Efficient VLMs with Visual Token Compression

Training-free

## Contextual Tokens Merging

- Merge the remaining tokens to avoid losing any small but potentially important information.



---

### Algorithm 2 Pseudocode for Contextual Tokens Merging.

---

```
# Remove dominant tokens
remaining = vanilla_tokens.mask(dominant_tokens)

# Split into target and merge tokens
# M represents the desired number of contextual tokens
targets, merge = uniform_split(remaining, M)

# Compute similarity based on the key values
similarity = bmm(to_merge.K, targets.K.transpose(1, 2))

# Assign each merge token to the most similar target
assign_idx = similarity.argmax(dim=2)

# Merge by averaging
context_tokens = avg_merge(assign_idx, targets, merge)
```

---

uniform\_split: Uniformly sample the target tokens, and the rest are the merge tokens; avg\_merge: Average merge the tokens based on the assigned indices.

# Efficient VLMs with Visual Token Compression

Training-free

## Experiments

Method	GQA	MMB	MME	POPE	SQA	VQA <sup>V2</sup>	VQA <sup>Text</sup>	MMMU	SEED	MMVet	LLaVA-B	Avg.
<i>Upper Bound, 576 Tokens (100%)</i>												
Vanilla (CVPR24)	61.9	64.7	1862	85.9	69.5	78.5	58.2	36.3	58.6	31.1	66.8	100%
	100%	100%	100%	100%	100%	100%	100%	100%	100%	100%	100%	100%
<i>Retain 192 Tokens (↓ 66.7%)</i>												
FastV (ECCV24)	52.7	61.2	1612	64.8	67.3	67.1	52.5	34.3	57.1	27.7	49.4	88.2%
	85.1%	94.6%	86.6%	75.4%	96.8%	85.5%	90.2%	94.5%	97.4%	89.7%	74.0%	
SparseVLM (2024.10)	57.6	62.5	1721	83.6	69.1	75.6	56.1	33.8	55.8	31.5	66.1	96.4%
	93.1%	96.6%	92.4%	97.3%	99.4%	96.3%	96.4%	93.1%	95.2%	101.3%	99.0%	
VisionZip	59.3	63.0	1782.6	85.3	68.9	76.8	57.3	36.6	56.4	31.7	67.7	98.5%
	95.8%	97.4%	95.7%	99.3%	99.1%	97.8%	98.5%	100.8%	96.2%	101.9%	101.3%	
VisionZip ‡	60.1	63.4	1834	84.9	68.2	77.4	57.8	36.2	57.1	32.6	66.7	99.1%
	97.1%	98.0%	98.5%	98.8%	98.1%	98.6%	99.3%	99.7%	97.4%	104.8%	99.9%	

‡ Fine-tuning visual projector; other frozen

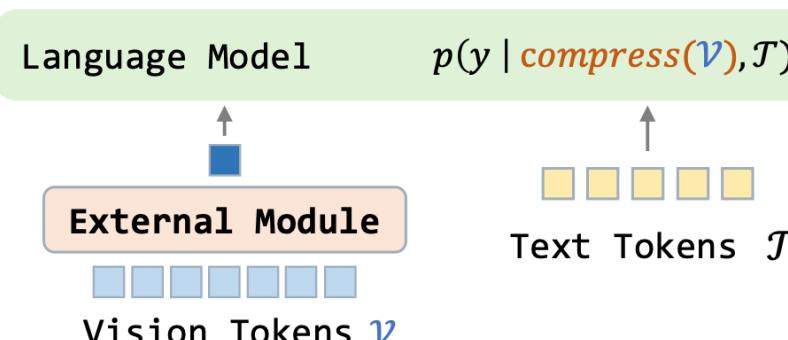
Retain 128 Tokens (↓ 77.8%)												
FastV (ECCV24)	49.6	56.1	1490	59.6	60.2	61.8	50.6	34.9	55.9	28.1	52.0	83.5%
	80.1%	86.7%	80.0%	69.4%	86.6%	78.7%	86.9%	96.1%	95.4%	90.9%	77.8%	
SparseVLM (2024.10)	56.0	60.0	1696	80.5	67.1	73.8	54.9	33.8	53.4	30	62.7	93.4%
	90.5%	92.7%	91.1%	93.7%	96.5%	94.0%	94.3%	93.1%	91.1%	96.5%	93.9%	
VisionZip	57.6	62.0	1761.7	83.2	68.9	75.6	56.8	37.9	54.9	32.6	64.8	97.6%
	93.1%	95.8%	94.6%	96.9%	99.1%	96.3%	97.6%	104.4%	93.7%	104.8%	97.6%	
VisionZip ‡	58.9	62.6	1823	83.7	68.3	76.6	57.0	37.3	55.8	32.9	64.8	98.4%
	95.2%	96.8%	97.9%	97.4%	98.3%	97.6%	97.9%	102.8%	95.2%	105.8%	97.0%	
Retain 64 Tokens (↓ 88.9%)												
FastV (ECCV24)	46.1	48.0	1256	48.0	51.1	55.0	47.8	34.0	51.9	25.8	46.1	75.6%
	74.5%	74.2%	67.5%	55.9%	73.5%	70.1%	82.1%	93.7%	88.6%	83.0%	69.0%	
SparseVLM (2024.10)	52.7	56.2	1505	75.1	62.2	68.2	51.8	32.7	51.1	23.3	57.5	85.8%
	85.1%	86.9%	80.8%	87.4%	89.4%	86.9%	89.0%	90.1%	87.2%	74.5%	86.1%	
VisionZip	55.1	60.1	1690	77.0	69.0	72.4	55.5	36.2	52.2	31.7	62.9	94.0%
	89.0%	92.9%	90.8%	89.6%	99.3%	92.2%	95.4%	99.7%	89.1%	101.9%	94.2%	
VisionZip ‡	57.0	61.5	1756	80.9	68.8	74.2	56.0	35.6	53.4	30.2	63.6	95.2%
	92.1%	95.1%	94.3%	94.2%	99.0%	94.5%	96.2%	98.1%	91.1%	97.1%	95.2%	

# Efficient VLMs with Visual Token Compression

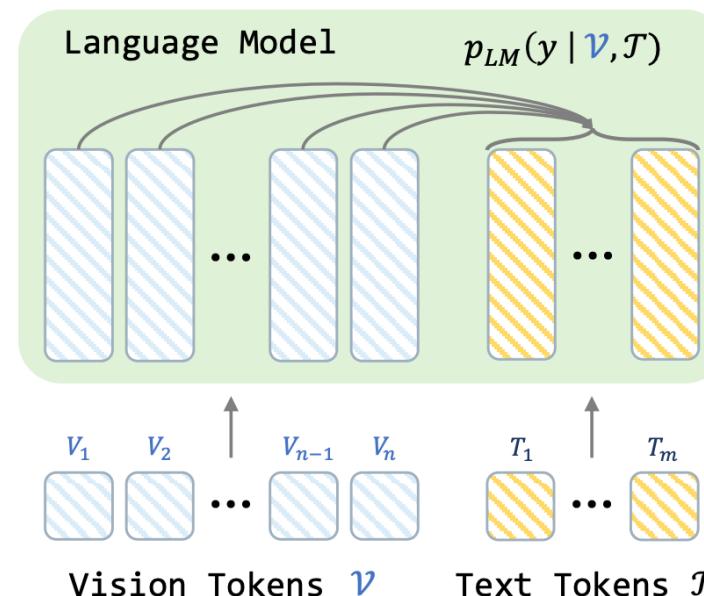
## VoCo-LLaMA

This work introduces a **learnable Vision Compression (VoCo) token** between visual and text tokens.

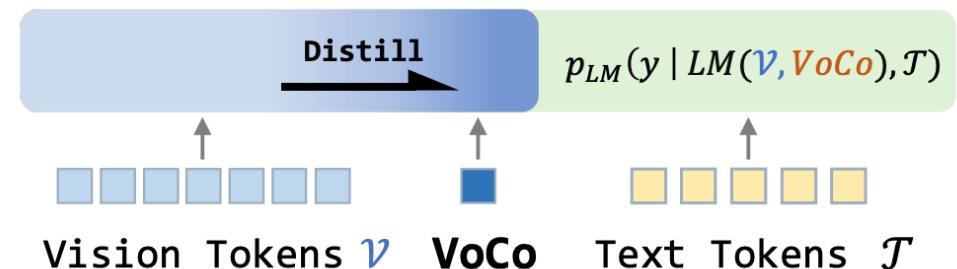
a) Previous methods



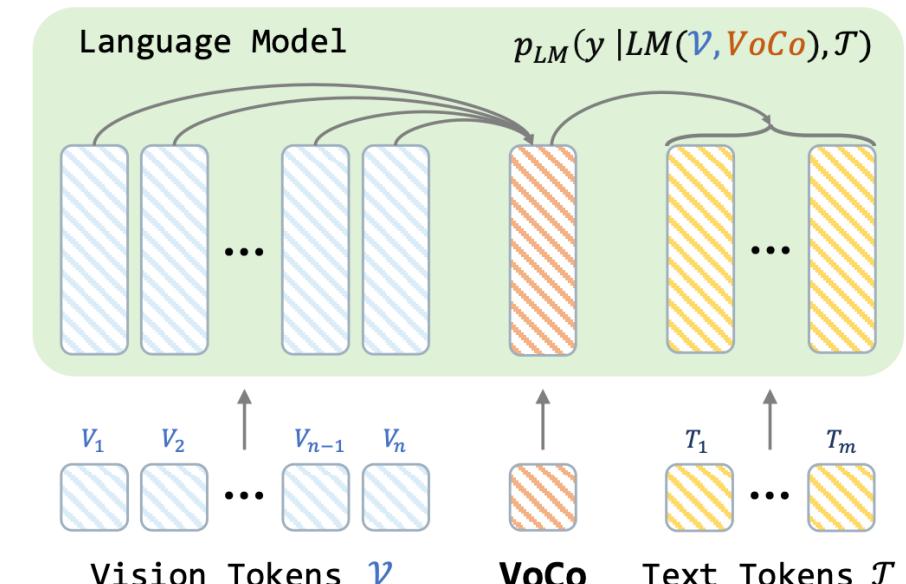
a) VLMs



b) VoCo-LLaMA



b) VoCo-LLaMA



VoCo-LLaMA: Towards Vision Compression with Large Language Models, in *CVPR2025*.

Learning to Compress Prompts with Gist Tokens, in *NeurIPS2023*.

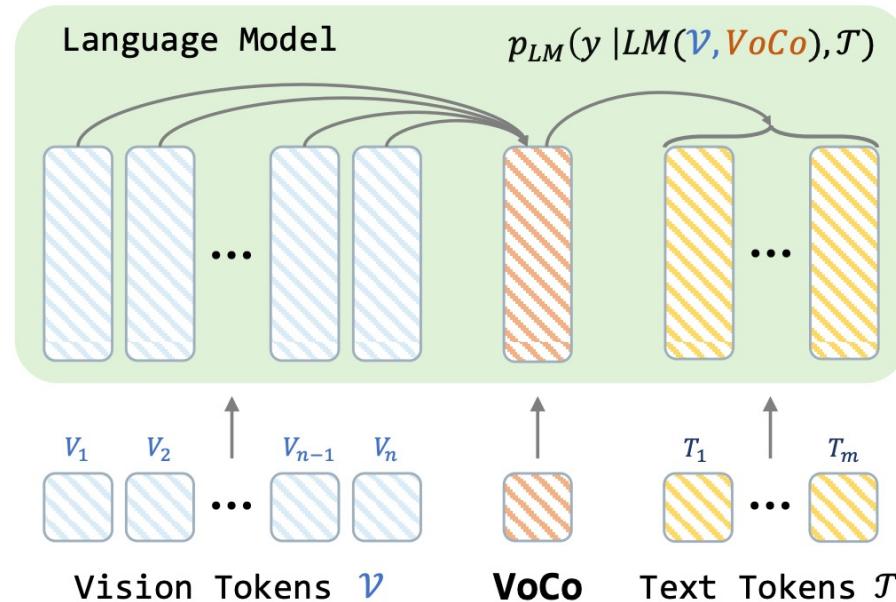
Tsinghua Uni. & Tencent

Stanford University

# Efficient VLMs with Visual Token Compression

Data-Driven Method

VoCo-LLaMA



576 $\times$  compression rate while maintaining 83.7% performance.

Modifying the attention mechanism, text tokens attend **solely** to VoCo tokens:

$$M_{ij} = \begin{cases} \text{True,} & \text{if } i \in \mathcal{T} \text{ and } j \in \text{VoCo}, \\ \text{False,} & \text{if } i \in \mathcal{T} \text{ and } j \in \mathcal{V}, \\ \text{True,} & \text{otherwise.} \end{cases}$$

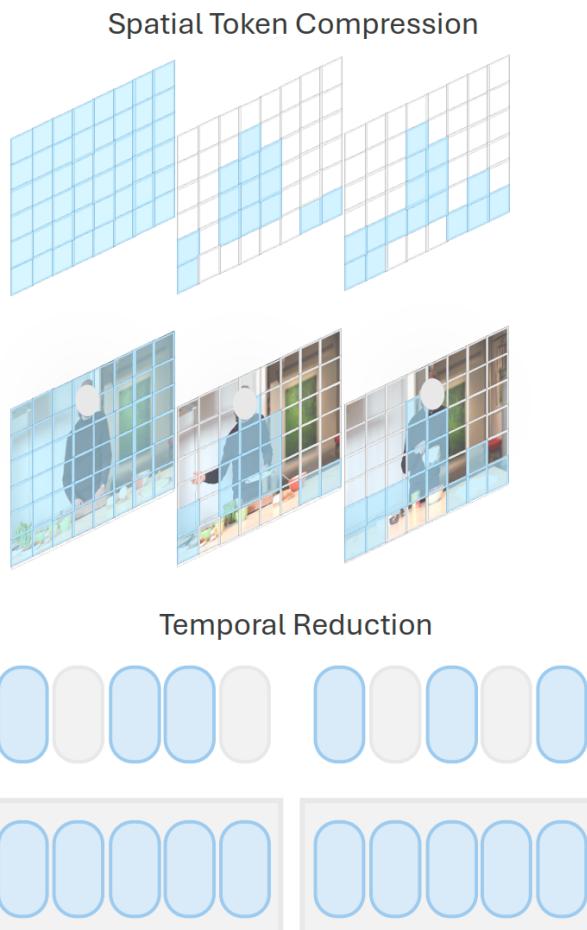
Distillation objective:

$$E_{\mathcal{V}, \mathcal{T}}[D_{KL}(p_{LM_o}(y | \mathcal{V}, \mathcal{T}) \| p_{LM_c}(y | \mathcal{C}, \mathcal{T}))]$$

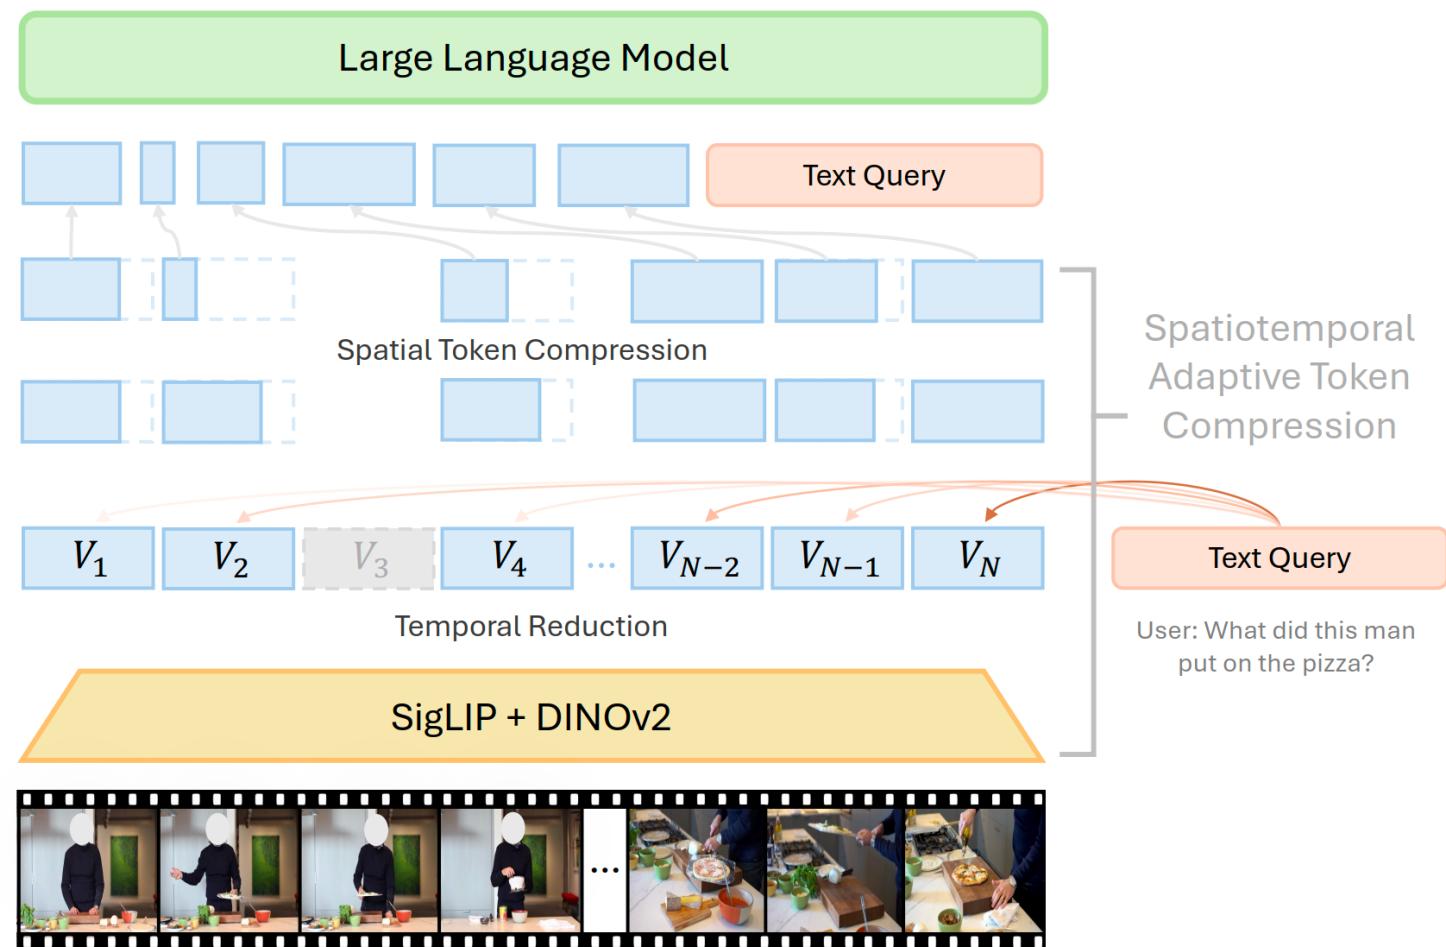
Token	MMB	GQA	VQA <sup>v2</sup>	SEED	Avg.
576	64.0	61.1	77.7	57.9	100%
128	<b>61.0</b>	59.8	<b>76.9</b>	<b>59.1</b>	<b>97.7%</b>
64	60.5	<b>60.4</b>	75.4	56.3	93.7%
32	59.4	60.2	75.3	56.2	92.6%
16	58.6	59.4	75.4	56.2	91.3%
8	58.7	59.2	75.3	56.3	91.3%
4	60.4	58.4	74.5	56.0	90.4%
2	60.1	57.7	73.5	55.0	87.8%
1	58.8	57.0	72.3	53.7	83.8%
1	22.3	37.7	41.2	36.9	0%

# Efficient VLMs with Visual Token Compression

LongVU



Model-driven  
Video method



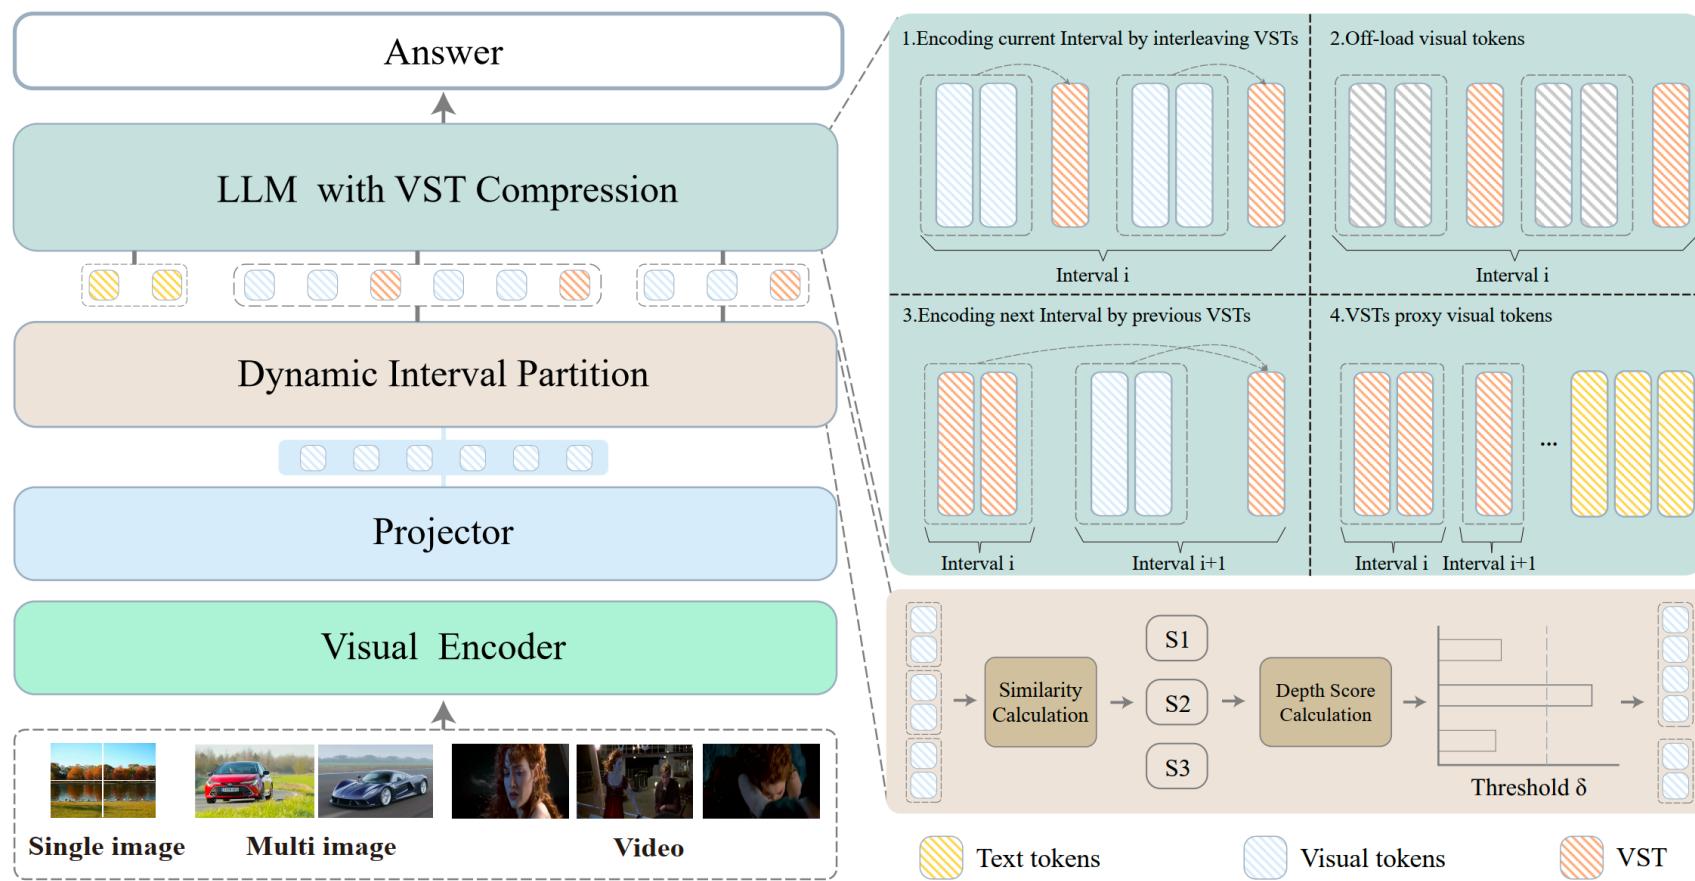
Step1: Temporal Reduction: DINOv2

Step2: Selective Feature Reduction via Cross-modal Query

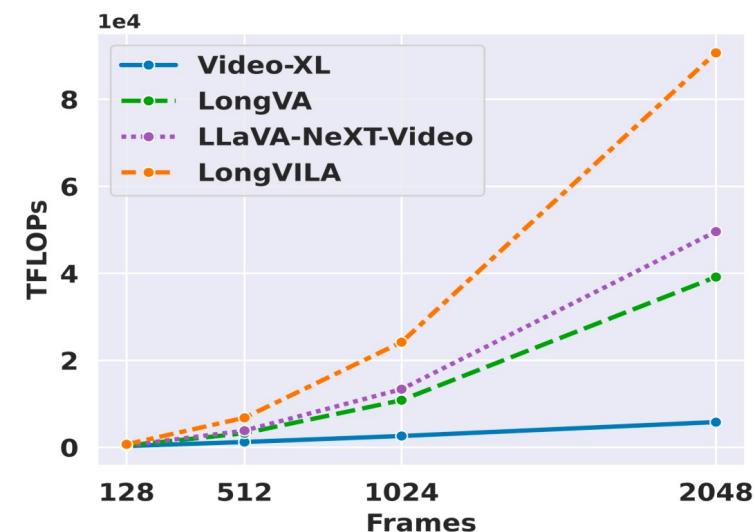
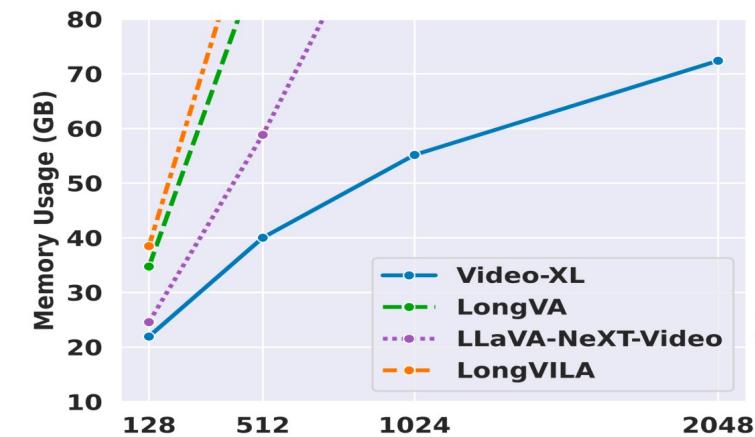
Step3: Spatial Token Compression (STC): pixel-level

# Efficient VLMs with Visual Token Compression

## VideoXL



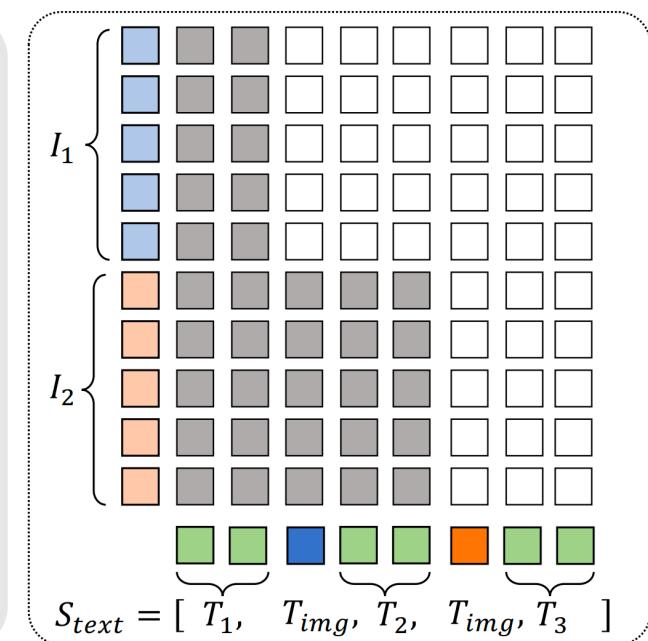
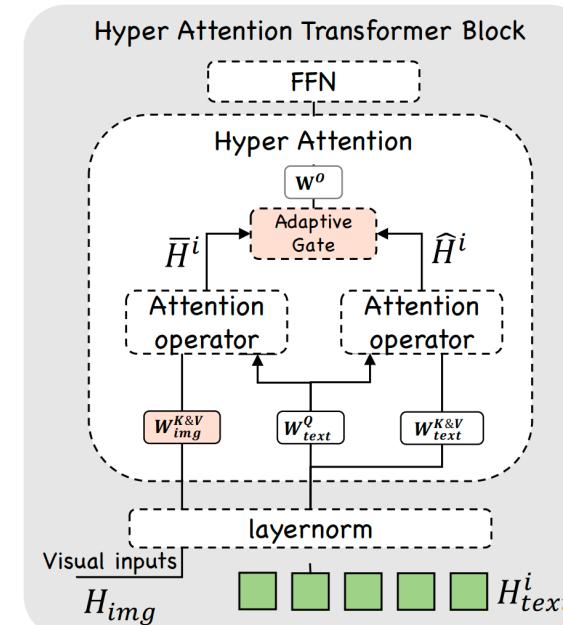
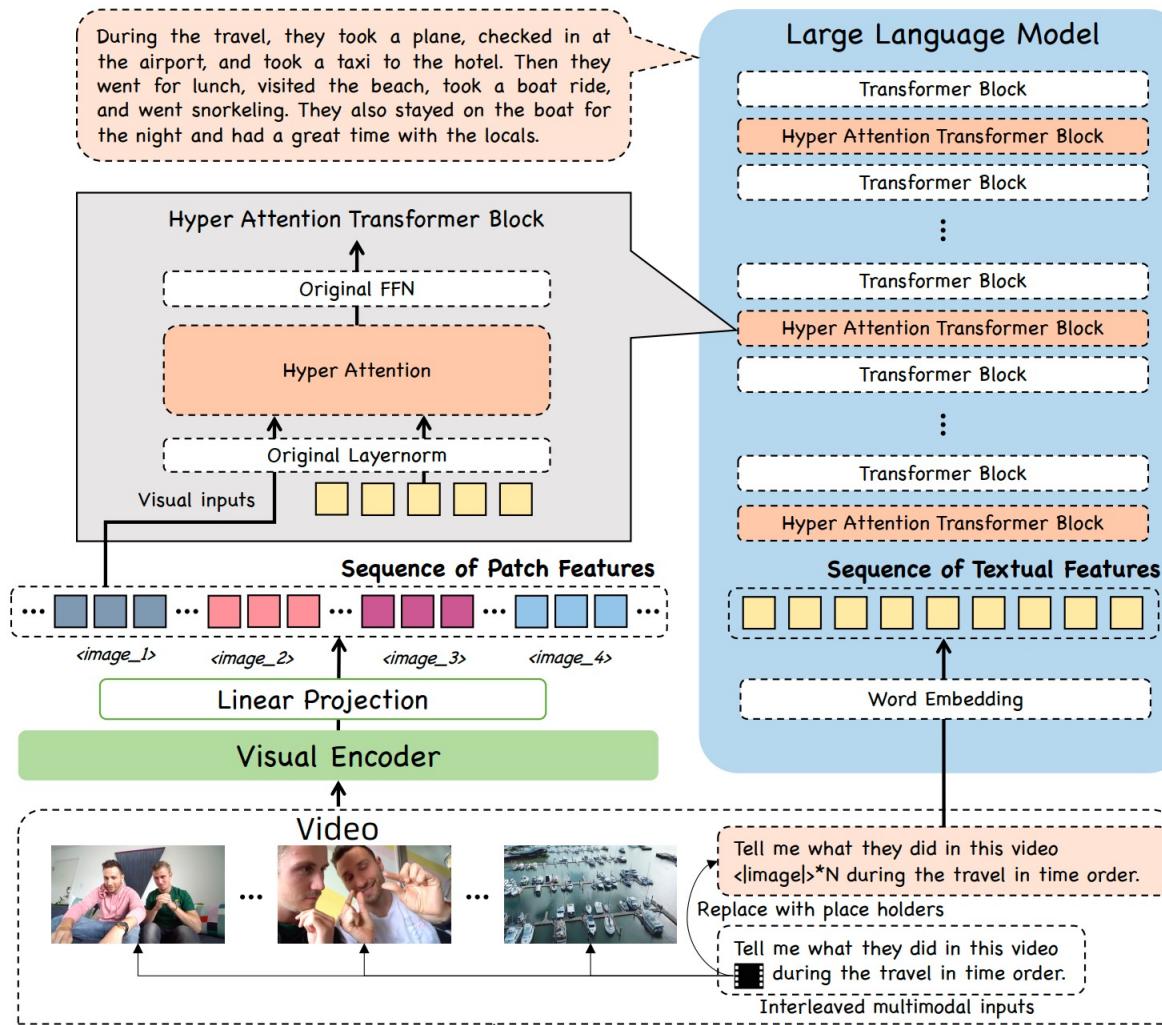
## Data-Driven Video Method



# Efficient VLMs with Visual Token Compression

Other paradigm

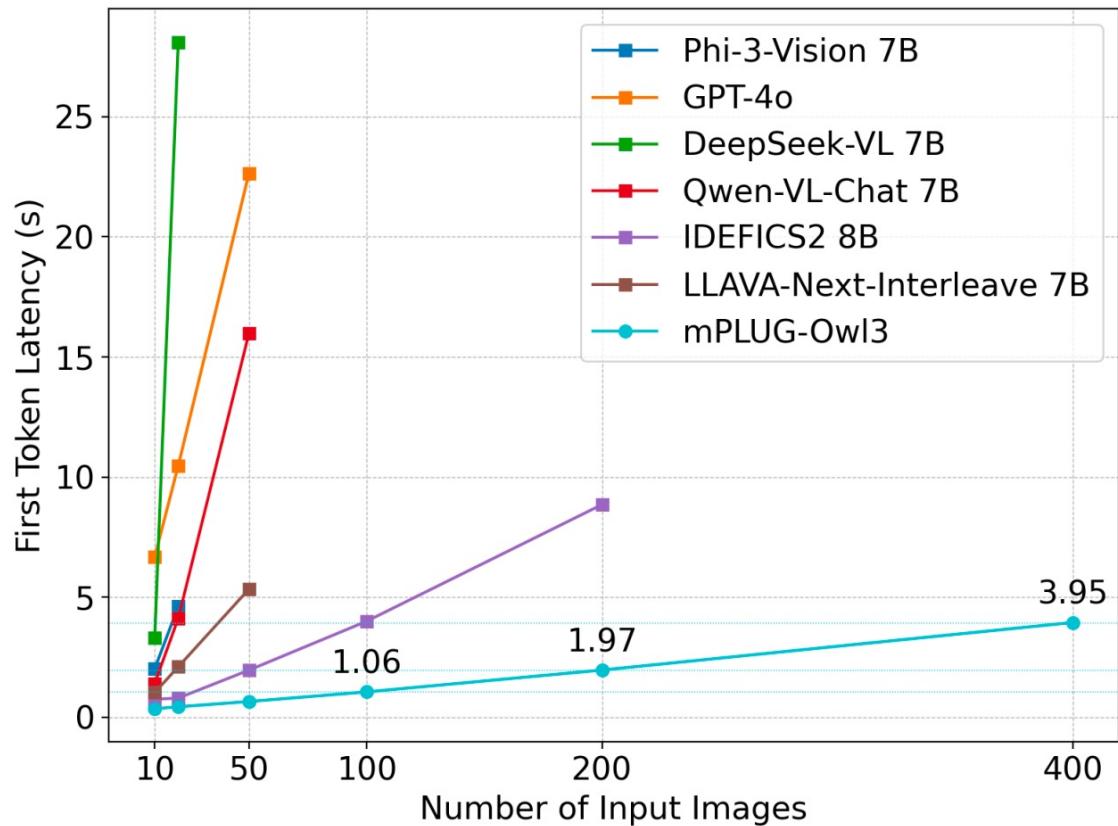
**mPLUG-Owl3: Only input text token and fuse visual tokens within attention block**



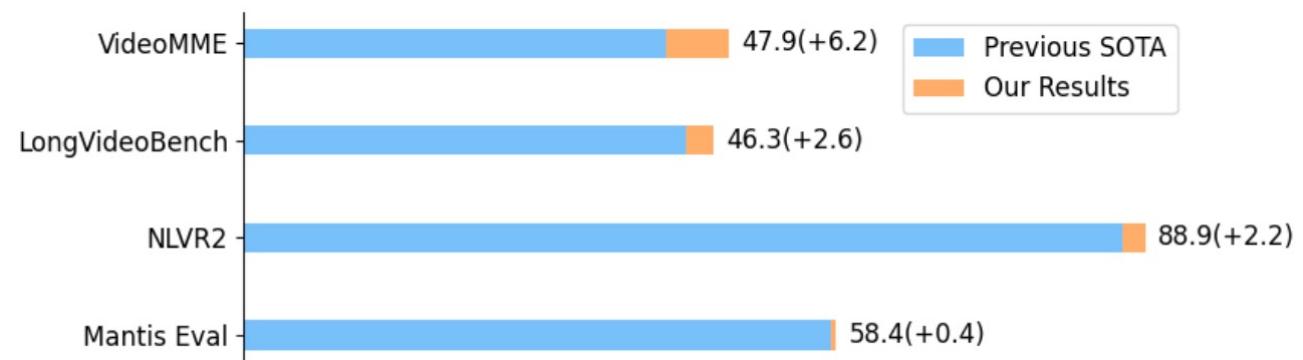
# Efficient VLMs with Visual Token Compression

Other paradigm

## Efficiency Comparisons



## Performance Comparisons



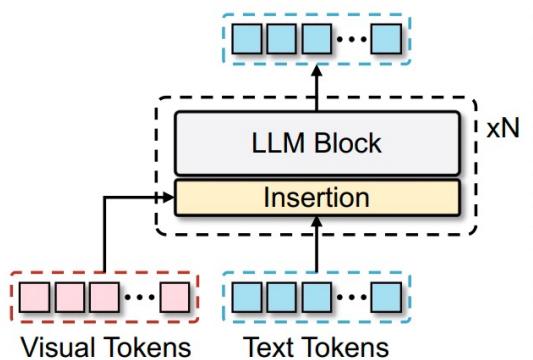
# Efficient VLMs with Visual Token Compression

Other paradigm

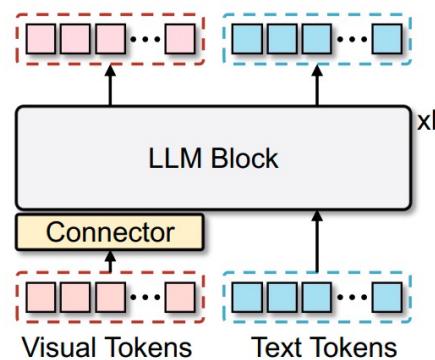
LaVi

Comparison with Current methods

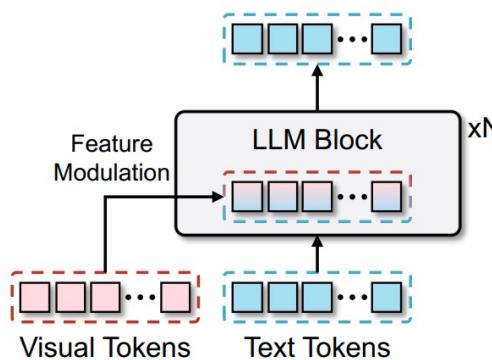
(a) Architectural Injection



(b) In-context Injection

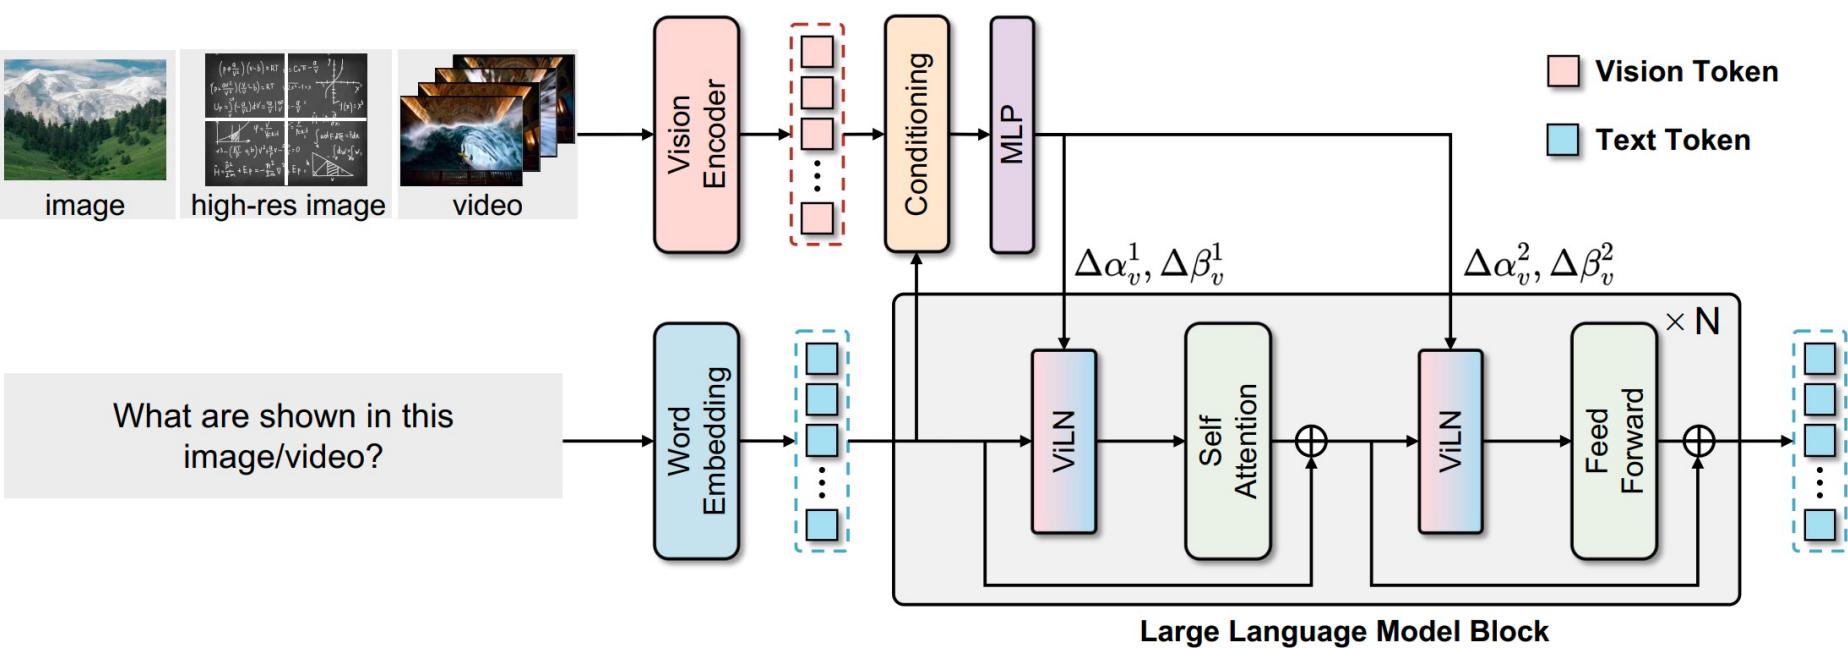


(c) Feature Modulation Injection



LaVi Framework

What are shown in this image/video?



# Efficient VLMs with Visual Token Compression

Other paradigm

## Feature Modulation Injection

Core insight: Vision-Infused Layer Normalization

Standard LN:

$$\text{LN}(t) = \alpha \odot \hat{t} + \beta \quad \alpha \text{ and } \beta \text{ are learnable affine parameters}$$

ViLN:

$$\text{ViLN}(t, v) = (\alpha + \Delta\alpha_v) \odot \hat{t} + (\beta + \Delta\beta_v)$$

$\Delta\alpha_v$  and  $\Delta\beta_v$  are *vision-conditioned deltas* generated from visual input  $v$ .

One before self-attention and One before FFN:

$$[\Delta\alpha_v^1, \Delta\beta_v^1, \Delta\alpha_v^2, \Delta\beta_v^2] = \text{Swish}(\text{Cond}(t, v)) \cdot W + b$$

Three Types of Conditioning Modules:

$$\text{Cond}_{mlp}(t_i, v) = \left[ \text{MLP}_{channel}\left( \left( \text{MLP}_{token}([t_i; v]^\top) \right)^\top \right) \right]_{t_i}$$

$$\text{Cond}_{conv}(t_i, v) = \left[ \text{Conv}_{point}\left( \sigma(\text{Conv}_{depth}([t_i; v])) \right) \right]_{t_i}$$

$$\text{Cond}_{attn}(t_i, v) = \text{Attention}(t_i \mathbf{W}_Q, v \mathbf{W}_K, v \mathbf{W}_V)$$

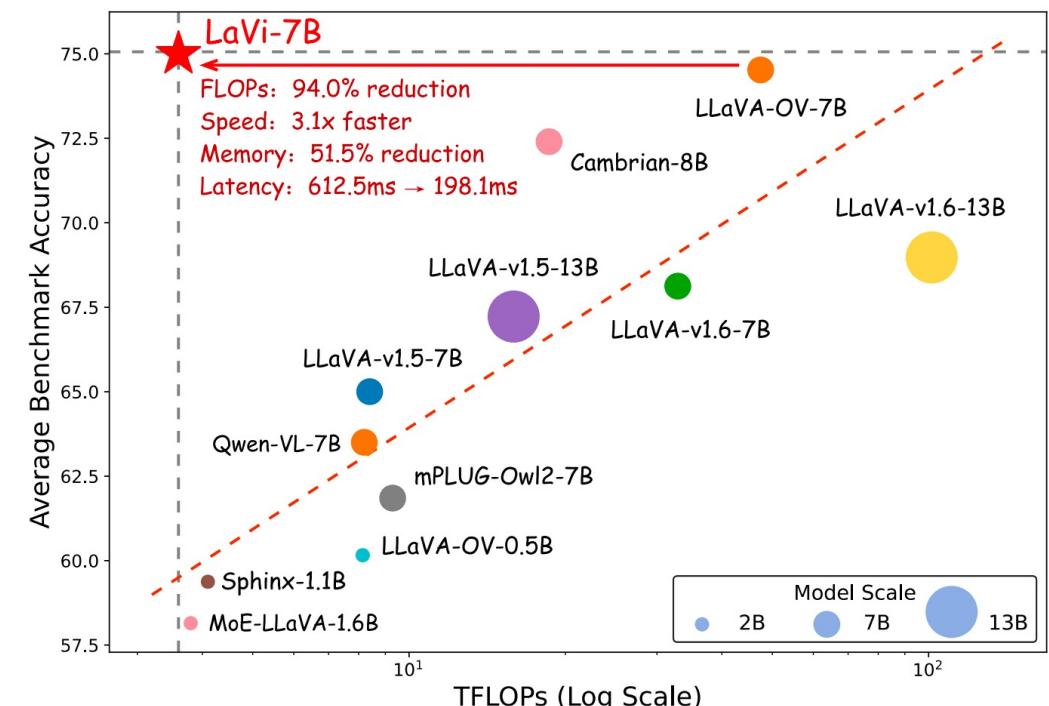
# Efficient VLMs with Visual Token Compression

Other paradigm

## Performance Comparisons

Method	LLM	Efficiency		Performance									
		FLOPs	Latency	VQA <sup>v2</sup>	GQA	VisWiz	SciQA	VQA <sup>T</sup>	POPE	MME <sup>P</sup>	MMB	SEED <sup>I</sup>	Avg.
<b>Baselines with <math>\leq 2B</math> parameters scale</b>													
MoE-LLaVA [36]	StableLM-1.6B	3.8	206.4	76.0	60.4	37.2	62.6	47.8	84.3	65.0	59.4	–	–
MobileVLM-V2 [18]	MLLaMA-1.4B	4.3	214.9	–	59.3	–	66.7	52.1	84.3	65.1	57.7	–	–
SPHINX-tiny [38]	TLLaMA-1.1B	4.1	212.3	74.7	58.0	49.2	21.5	57.8	82.2	63.1	56.6	25.2	54.3
LLaVA-OV [27]	Qwen2-0.5B	7.8	228.0	78.5	58.0	51.4	67.2	65.9	86.0	61.9	52.1	65.5	65.2
<b>Baselines with <math>\leq 8B</math> parameters scale</b>													
Qwen-VL-Chat [8]	Qwen-7B	8.2	239.4	78.2	57.5	38.9	68.2	61.5	–	74.4	60.6	65.4	–
mPLUG-Owl2 [65]	LLaMA2-7B	9.3	278.6	79.4	56.1	54.5	68.7	54.3	–	72.5	64.5	57.8	–
Cambrian-1 [57]	LLaMA3-8B	18.6	393.7	–	64.6	–	80.4	71.7	–	77.4	75.9	74.7	–
LLaVA-v1.5 [39]	Vicuna-7B	8.4	254.4	78.5	62.0	50.0	66.8	58.2	85.9	75.5	64.3	66.1	67.5
LLaVA-v1.6 [40]	Vicuna-7B	32.9	502.4	81.8	64.2	57.6	70.1	64.9	86.5	76.0	67.4	70.2	71.0
LLaVA-OV [27]	Qwen2-7B	60.4	612.5	84.5	62.2	53.0	96.0	76.1	87.4	79.0	80.8	75.4	77.2
<b>Ours</b>													
LaVi-Image	Vicuna-7B	0.6	110.8	79.6	63.0	52.9	67.8	58.4	86.9	75.2	64.8	67.5	68.5
△ compare to LLaVA-v1.5		<b>7.1%</b>	<b>43.6%</b>	<b>+1.1</b>	<b>+1.0</b>	<b>+2.9</b>	<b>+1.0</b>	<b>+0.2</b>	<b>+1.0</b>	<b>-0.3</b>	<b>+0.5</b>	<b>+1.4</b>	<b>+1.0</b>
LaVi-Image (HD)	Vicuna-7B	1.7	148.6	81.4	63.7	57.8	71.7	64.3	87.0	77.5	68.1	71.6	71.5
△ compare to LLaVA-v1.6		<b>5.2%</b>	<b>29.6%</b>	<b>-0.4</b>	<b>-0.5</b>	<b>+0.2</b>	<b>+1.6</b>	<b>-0.6</b>	<b>+0.5</b>	<b>+1.5</b>	<b>+0.7</b>	<b>+1.4</b>	<b>+0.5</b>
LaVi	Qwen2-7B	3.6	198.1	84.0	65.0	53.8	95.4	77.0	87.1	80.9	79.3	76.9	77.7
△ compare to LLaVA-OV		<b>6.0%</b>	<b>32.3%</b>	<b>-0.5</b>	<b>+2.8</b>	<b>+0.8</b>	<b>-0.6</b>	<b>+0.9</b>	<b>-0.3</b>	<b>+1.9</b>	<b>-1.5</b>	<b>+1.5</b>	<b>+0.5</b>

## Efficiency Comparisons



Without comparison with mPLUG-Owl3

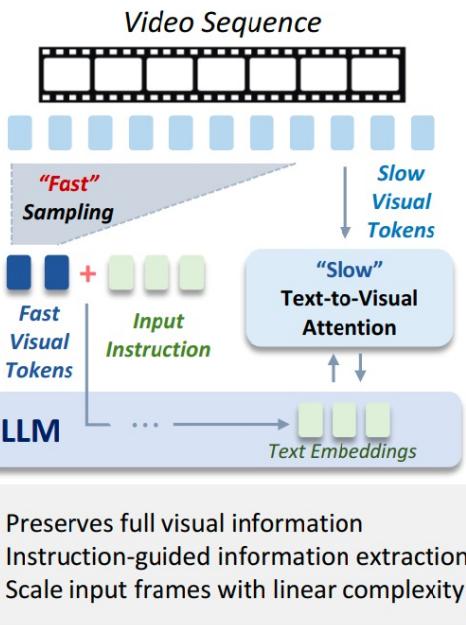
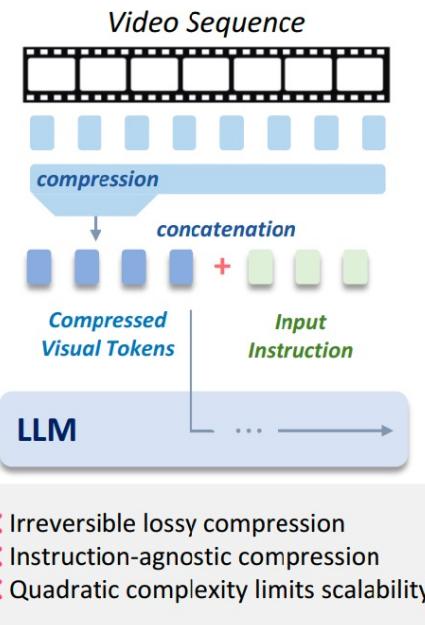
IACAS

# Efficient VLMs with Visual Token Compression

## Slow-fast MLLM

Other paradigm  
Video method

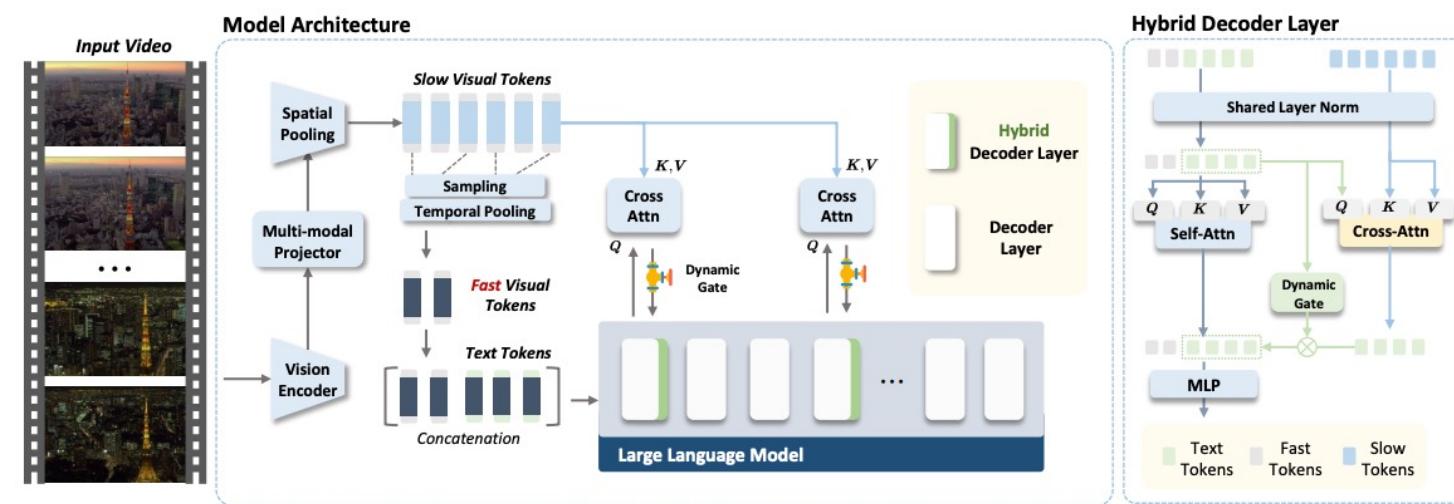
Comparison with Current method



Mainstream Architecture

Our Slow-Fast Architecture

Framework



# Efficient VLMs with Visual Token Compression

## Conclusion and Future direction

- **Model-driven Approaches**

Numerous recent studies have emerged, though the potential for further improvement is becoming limited—particularly for image-based VLMs.

- **Data-driven Approaches**

Demonstrate significant advantages when dealing with extremely fewer visual tokens;

Develop large-scale token ranking datasets;

Propose methods with strong generalization capabilities.

- **Other Paradigms**

Develop more effective **Vision-Infused Modules**;

Research in this area remains limited, especially for Video-LLMs.

*Thanks !*

Wentong LI (李文通)

Homepage: <https://cslwt.github.io/>

Wechat: 17795837723